



# Randomized Composable Core-sets for Distributed Optimization

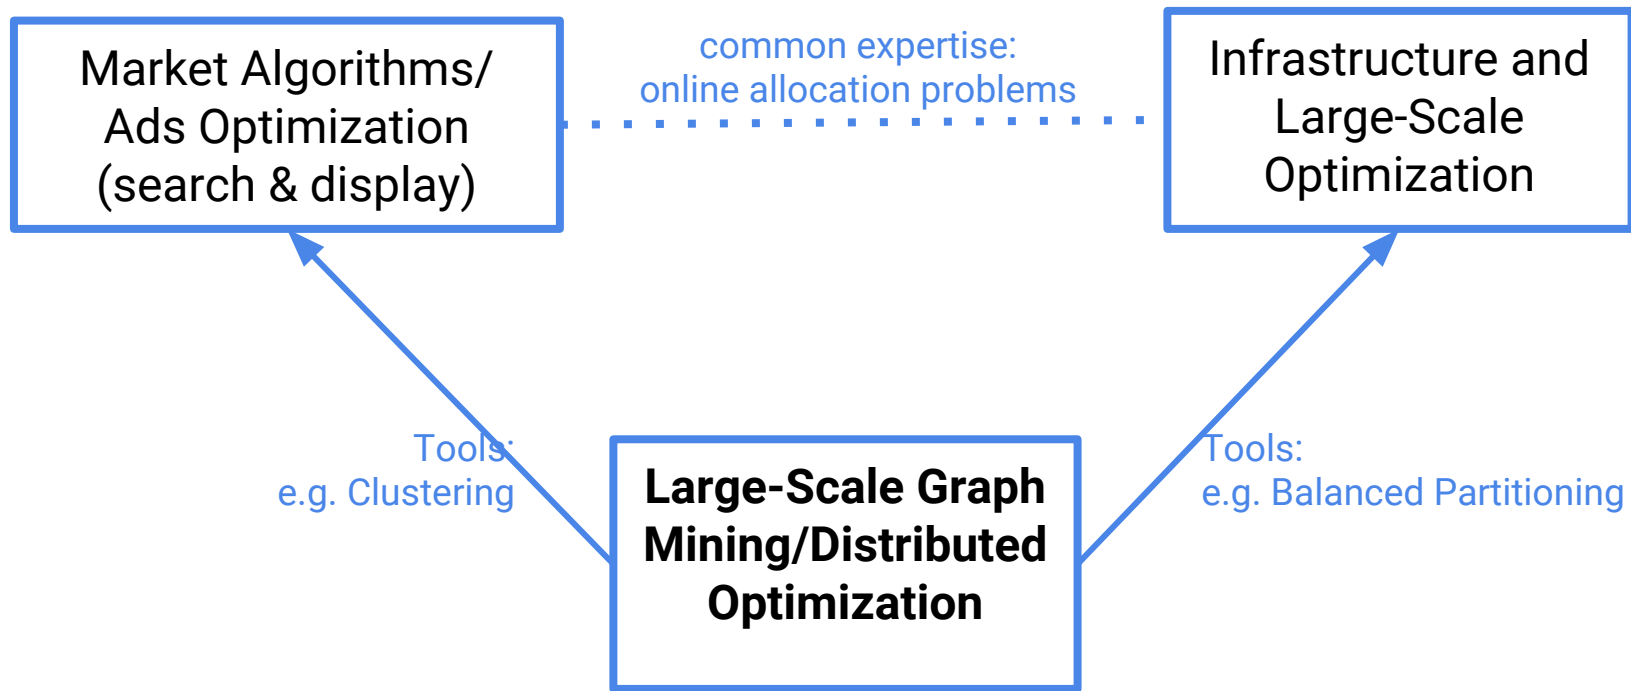
Vahab Mirrokni

Algorithms Research Group,  
Google Research, New York

*Mainly based on joint work with:*

Hossein Bateni, Aditya Bhaskara,  
Hossein Esfandiari, Silvio Lattanzi,  
Morteza Zadimoghaddam

# Our team: Google NYC Algorithms Research Teams



# Three most popular techniques applied in our tools


1. Local Algorithms: Message Passing/Label Propagation/Local Random Walks
  - e.g., similarity ranking via PPR etc, Connected Components
  - Connected components code that's 10-50 times faster the state-of-the-art
2. Embedding/Hashing/Sketching Techniques
  - e.g., linear embedding for balanced graph partitioning to minimize cut
  - Improves the state-of-the-art by 26%. Improved flash bandwidth for search backend by 25%. Paper appeared in WSDM'16.
3. Randomized Composable Core-sets for Distributed Computation: **This Talk**

# Agenda



- **Composable core-sets: Definitions & Applications**
  - Applications in Distributed & Streaming settings
  - Applications: Feature Selection, Diversity in Search & Recom.
- **Composable Core-sets for Four Problems: Survey**
  - Diversity Maximization(PODS'14, AAI'17),  
Clustering(NIPS'14), **Submodular Maximization**(STOC'15),  
and Column Subset Selection (ICML'16)
- **Sketching for Coverage Problems (on arXiv)**
  - Sketching Technique

# Composable Core-Sets for Distributed Optimization



# Composable Core-sets

Setup: Consider partitioning data set  $T$  of elements into  $m$  sets  $(T_1, T_2, \dots, T_m)$ .

$$T = T_1 \cup T_2 \cup \dots \cup T_m$$

Goal: Given a set function  $f$ , find a subset  $S^*$  with  $|S^*| \leq k$ , optimizing  $f(S^*)$ .

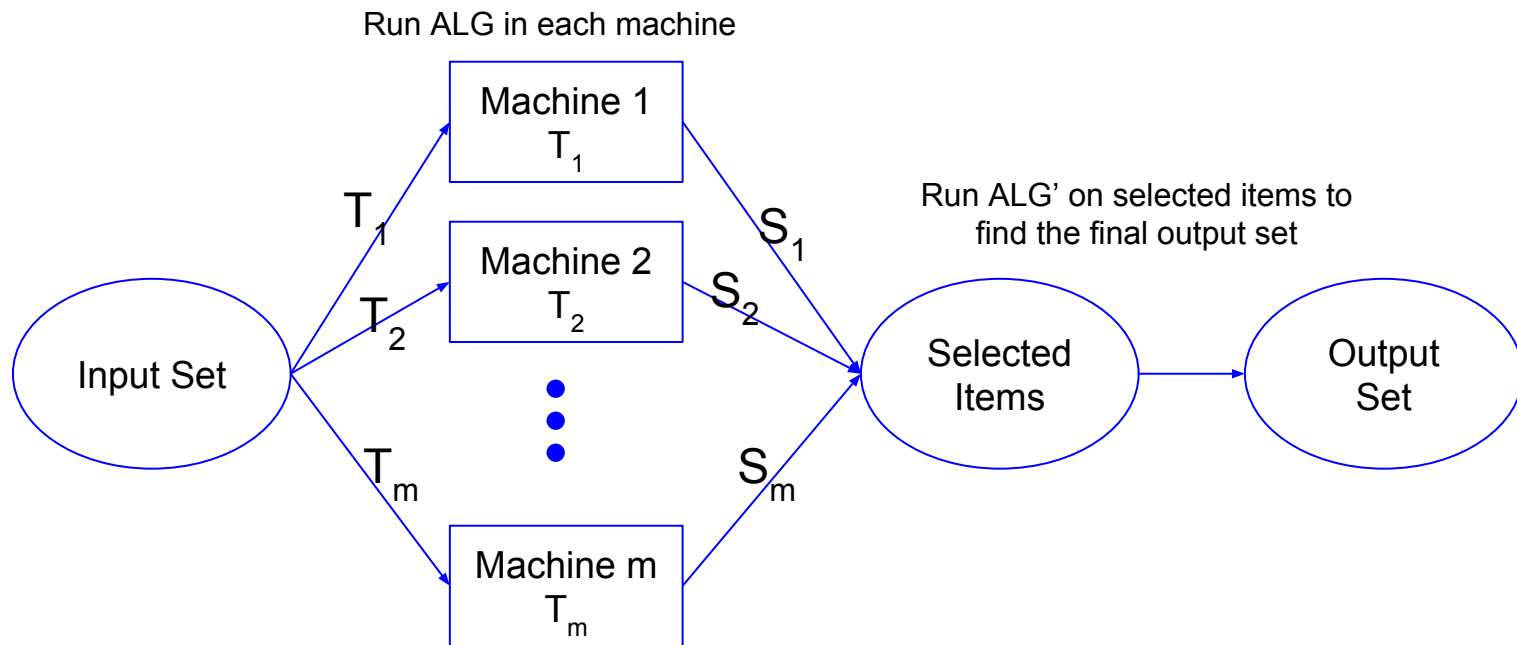
$$\text{opt}(T) = f(S^*)$$

Find: *small* core-set  $S_1 \subseteq T_1, S_2 \subseteq T_2, \dots, S_m \subseteq T_m$  such that

*optimum solution in union of core-sets approximates the optimum solution of  $T$*

$$\frac{1}{c} \text{opt}(S_1 \cup S_2 \dots \cup S_m) \leq \text{opt}(T_1 \cup T_2 \dots \cup T_m) \leq c \times \text{opt}(S_1 \cup S_2 \dots \cup S_m)$$


# Application in MapReduce/Distributed Computation



E.g., two rounds of MapReduce

# Application in Streaming Computation

- **Streaming Computation:**
  - Processing sequence of  $n$  data points “on the fly”
  - Limited storage
- Use C-composable core-set of size  $k$ , for example:
  - Chunks of size  $\sqrt{nk}$ , thus number of chunks is  $\sqrt{n/k}$
  - Compute core-set of size  $k$  for each chunk
  - Total space:  $k\sqrt{n/k} + \sqrt{nk} = O(\sqrt{nk})$





# Overview of recent theoretical results

Need to solve (combinatorial) optimization problems on large data

## 1. **Diversity Maximization,**

- **PODS'14** by IndykMahdianMahabadiMirrokni
- for Feature Selection in **AAAI'17** by AbbasiGhadiriMirrokniZadimoghaddam

## 2. **Capacitated $\ell_p$ Clustering, NIPS'14** by BateniBhaskaraLattanziMirrokni

## 3. **Submodular Maximization, STOC'15** by MirrokniZadimoghaddam

## 4. **Column Subset Selection** (Feature Selection), **ICML'16** by Alschulter et al.

## 5. Coverage Problems: **Submitted** by BateniEsfandiariMirrokni



# Applications: Diversity & Submodular Maximization

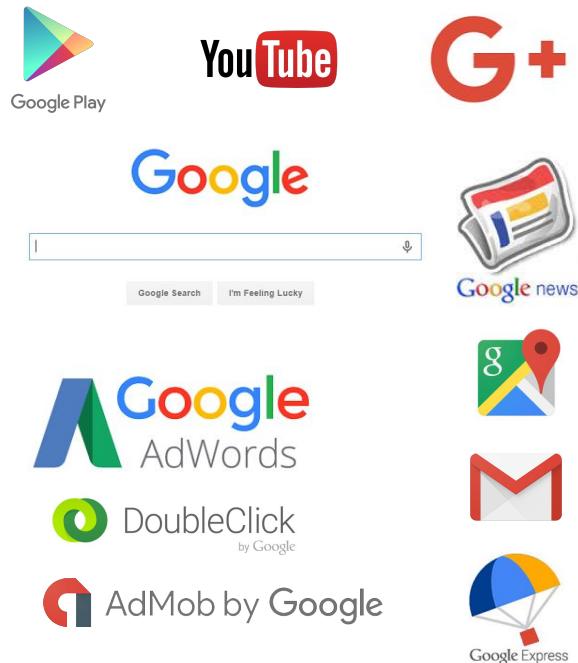
## Diverse suggestions

- Play apps
- Campaign keywords
- Search results
- News articles
- YouTube videos

## Data summarization

- Feature selection

## Exemplar sampling





# Feature selection

We have

- Data points (docs, web pages, etc.)
- Features (topics, etc.)

**Goal:** pick a small set of “representative” features



Emotion



Hotel



Weather



Movie



Gaming



Smartphone



Car



World



Finance



Hospital



Cloud



Theatre



Software



Laptop



Boat



Home



Money



Business



Laundry



Camera



Education



Security



Train



Shopping

# Five Problems Considered

**General:** Find a set  $S$  of  $k$  items & maximize/minimize  $f(S)$ .

- **Diversity Maximization:** Find a set  $S$  of  $k$  points, and maximize the sum of pairwise distances i.e.  $\max \text{diversity}(S) = \sum_{i,j \in S} \text{dist}(i, j)$ .
- **Capacitated/Balanced Clustering:** Find a set  $S$  of  $k$  centers and cluster nodes around them while minimizing the sum of distances to  $S$ .
- **Coverage/Submodular Maximization:** Find a set  $S$  of  $k$  items. Maximize submodular function  $f(S)$ . Generalizing set cover.
- **Column subset selection:** Given a matrix  $A$ , find a set  $S$  of  $k$  columns.
  - Minimize  $\|A - \Pi_{A[S]} A\|_{\mathcal{F}}^2$

# Diversity Maximization Problem

- Given: A set of  $n$  points in a metric space  $(X, \text{dist})$
- Find a set  $S$  of  $k$  points
- Goal: maximize  $\text{diversity}(S)$  i.e.

**$\text{diversity}(S)$  = sum of pairwise distances of points in  $S$ .**

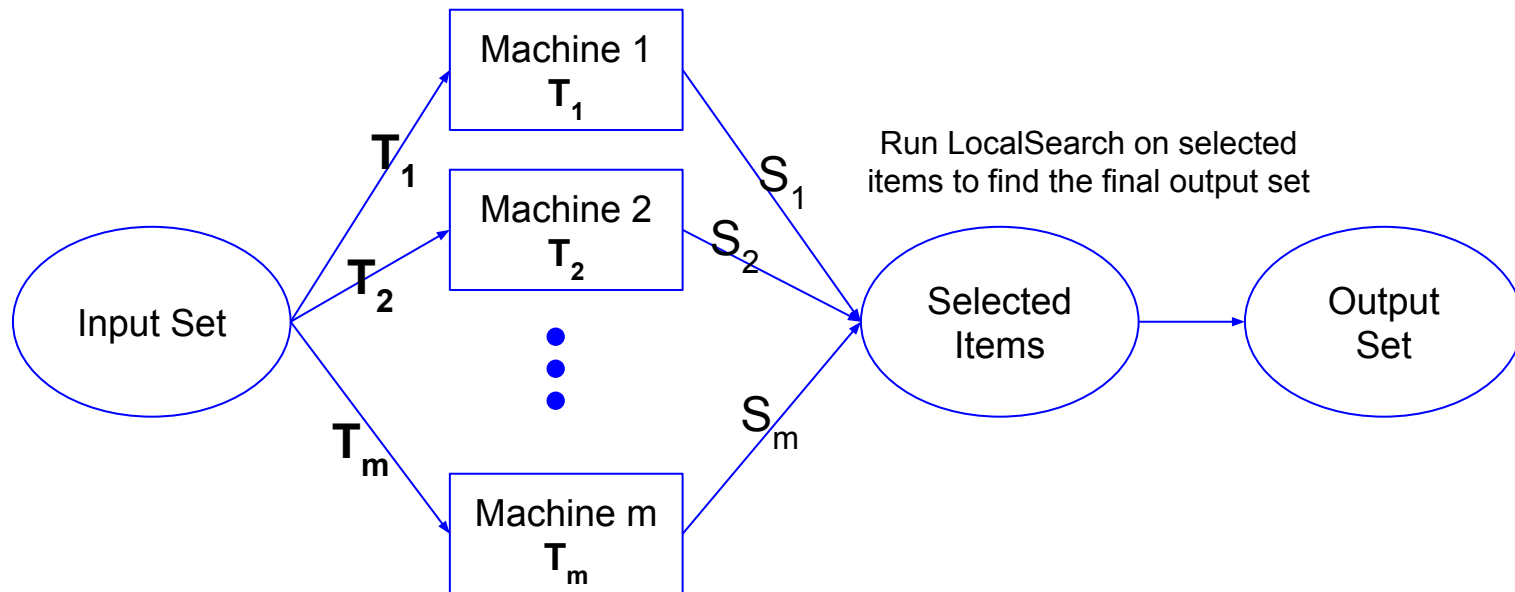
$$\text{diversity}(S) = \sum_{i,j \in S} \text{dist}(i, j)$$

- Background: Max Dispersion (Halldorson et al, Abbassi et al)
- Useful for feature selection, diverse candidate selection in Search, representative centers...

# Core-sets for Diversity Maximization

## Two rounds of MapReduce

Run LocalSearch on each machine




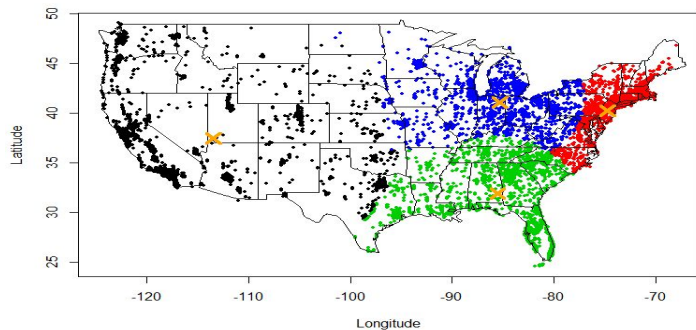
- Arbitrary Partitioning works. Random partitioning is better.

# Composable Core-set Results for Diversity Maximization

- Theorem(IndykMahabadiMahdianM.'14): The local search algorithm computes a *constant-factor* composable core-set for maximizing *sum of pairwise distances* in **2 rounds**:
- Theorem(EpastoM.ZadiMoghaddam'16): A sampling+greedy algorithm computes a randomized **2-approximate** composable *small-size core-set* for diversity maximization in **one round**.
  - *randomized*: works under random partitioning
  - *small-size*: size of core-set is less than  $k$ .

# Distributed Clustering Problems

**Clustering:** Divide data into groups containing “nearby” points



**Minimize:**

**k-center :**  $\max_i \max_{u \in S_i} d(u, c_i)$

**k-means :**  $\sum_i \sum_{u \in S_i} d(u, c_i)^2$

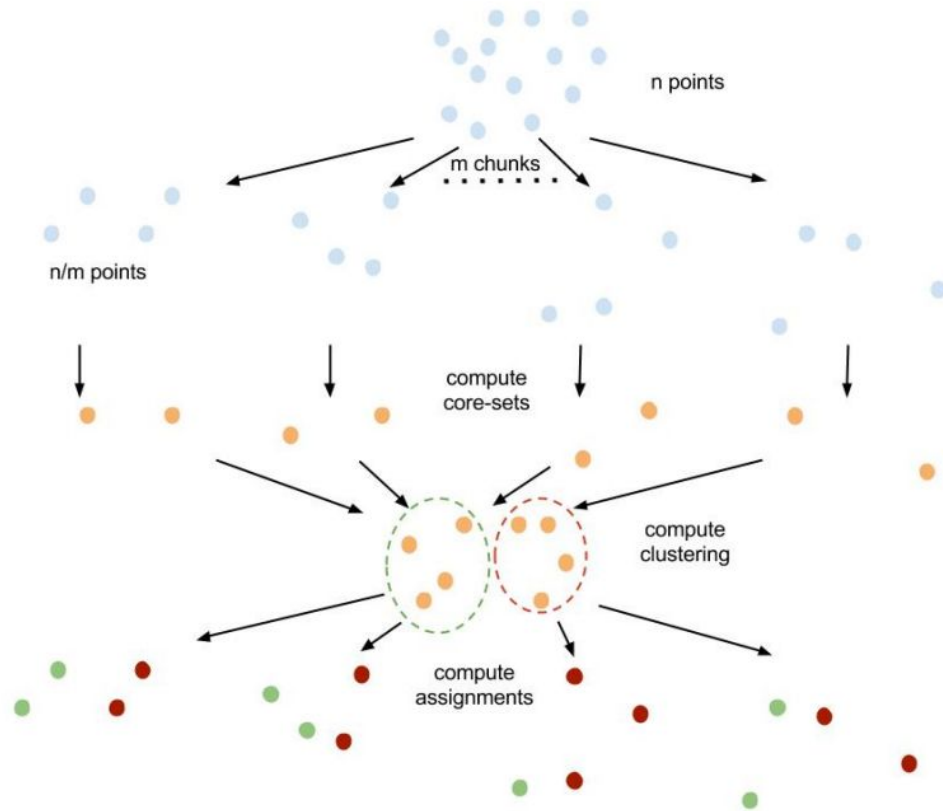
**k-median :**  $\sum_i \sum_{u \in S_i} d(u, c_i)$

Metric space  $(d, X)$

$\alpha$ -approximation  
algorithm: cost less than  
 $\alpha * \text{OPT}$



# Mapping Core-sets for Capacitated Clustering



# Capacitated $\ell_p$ clustering

**Problem:** Given  $n$  points in a metric space, find  $k$  centers and assign points to centers, *respecting capacities*, to minimize  $\ell_p$  norm of the distance vector.

- Generalizes *balanced  $k$ -median*,  *$k$ -means* &  *$k$ -center*.
- Objective is *not* minimizing cut size (cf. “balanced partitioning” in the library)

**Theorem:** For any  $p$  and  $k \leq \sqrt{n}$ , distributed balanced clustering with

- **approx ratio:** ‘small constant’ \* ‘best single machine guarantee’
  - **# rounds:** 2
  - **memory:**  $(n/m)^2$  with  $m$  machines
- Improves [BMVKV’12] and [BEL’13]

(Bateni, Bhaskara, Lattanzi, Mirrokni, NIPS’14)

# Empirical study for distributed clustering

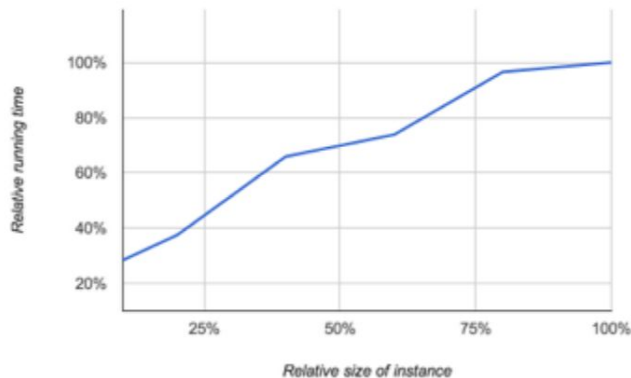
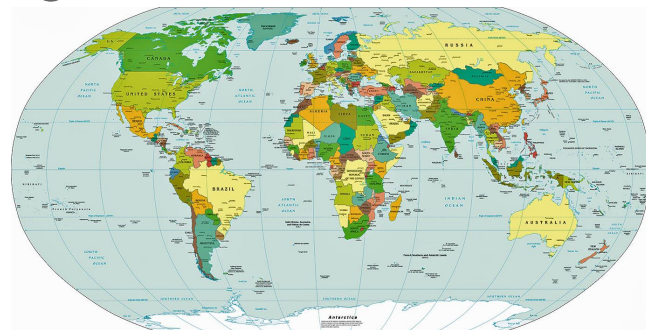
Test in terms of **scalability** and **quality of solution**

Two “base” instances & subsamples

- US graph ~30M nodes
- World graph ~500M nodes

	Size of seq. inst	Increase in OPT
US	1/300	1.52
World	1/1000	1.58

**Quality:** pessimistic analysis



**Sublinear** running time **scaling**

# Submodular maximization

**Problem**: Given  $k$  & submodular function  $f$ , find set  $S$  of size  $k$  that maximizes  $f(S)$ .

Some applications

- Data summarization
- Feature selection
- Exemplar clustering

**Special case**: “coverage maximization”: Given a family of subsets, choose a subfamily of  $k$  sets, and maximize cardinality of union.

- cover various topics/meanings
- target all kinds of users

# Submodular maximization

**Problem**: Given  $k$  & submodular function  $f$ , find set  $S$  of size  $k$  that maximizes  $f(S)$ .

Some applications

- Data summarization
- Feature selection
- Exemplar clustering

**Special case**: “coverage maximization”: Given a family of subsets, choose a subfamily of  $k$  sets, and maximize cardinality of union.

- cover various topics/meanings
- target all kinds of users

[IMMM'14] Bad News: **No deterministic** composable core-set with approx  $\leq \frac{\sqrt{k}}{\log k}$

# Submodular maximization

**Problem**: Given  $k$  & submodular function  $f$ , find set  $S$  of size  $k$  that maximizes  $f(S)$ .

Some applications

- Data summarization
- Feature selection
- Exemplar clustering

**Special case**: “coverage maximization”: Given a family of subsets, choose a subfamily of  $k$  sets, and maximize cardinality of union.

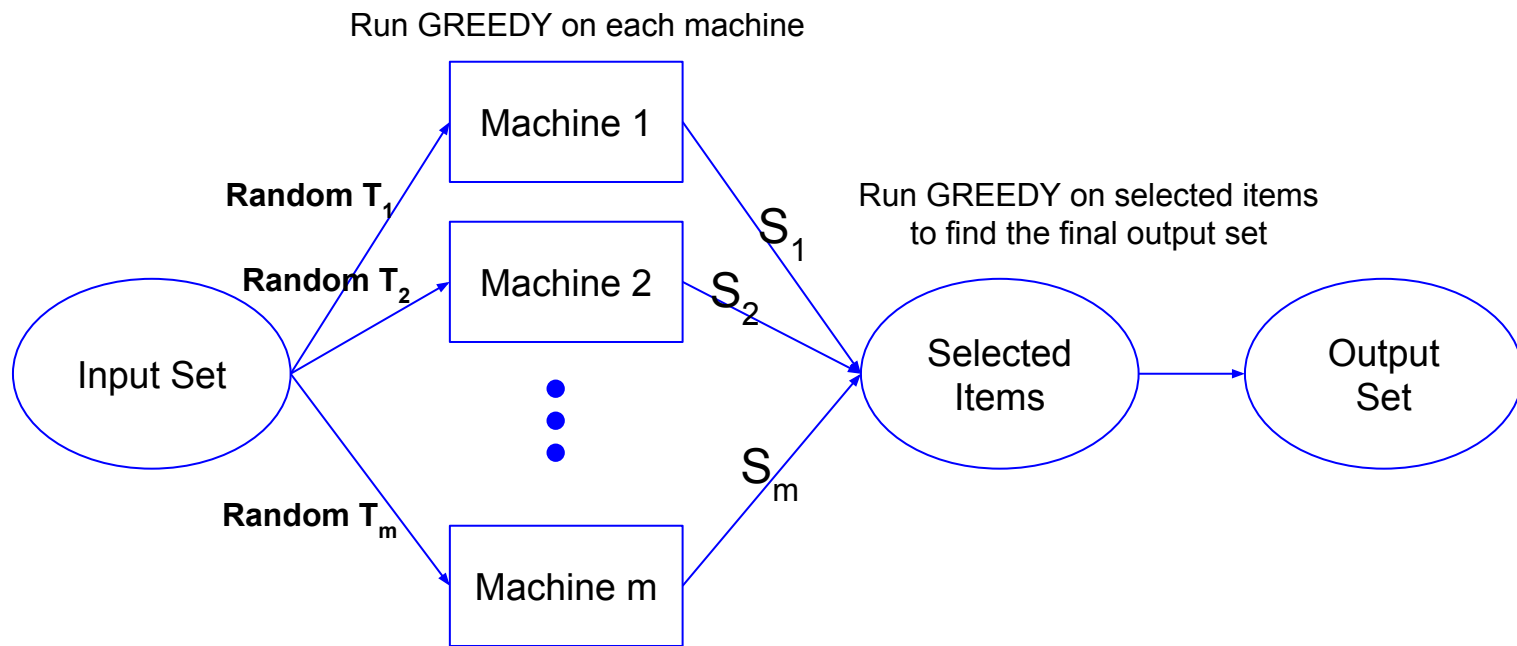
- cover various topics/meanings
- target all kinds of users

[IMMM'14] Bad News: No **deterministic** composable core-set with approx  $\leq \frac{\sqrt{k}}{\log k}$

**Randomization** is necessary and useful:

- Send each set **randomly** to some machine
- Build a coreset on each machine by greedy algorithm

# Randomization to the Rescue: Randomized Core-sets



Two rounds of MapReduce

# Results for Submodular Maximization: MZ (STOC'15)

- A class of **0.33**-approximate randomized composable core-sets of size **k** for **non-monotone** submodular maximization. For example, Greedy Algorithm.
- **Hard** to go beyond  $\frac{1}{2}$  approximation with size **k**. **Impossible** to get better than  $1-1/e$ .
- **0.58**-approximate randomized composable core-set of **size  $4k$**  for monotone  $f$ . Results in **0.54**-approximate distributed algorithm in two rounds with linear communication complexity.
- For **small-size composable core-sets** of **k'** less than **k**:  **$\sqrt{k'/k}$** -approximate randomized composable core-set.



# Low-Rank Approximation

Given (large) matrix  $A$  in  $\mathbb{R}^{m \times n}$  and target rank  $k \ll m, n$ :

$$\arg \min_{X, \text{rank}(X)=k} \|A - X\|_F^2$$

- Optimal solution: k-rank SVD
- Applications:
  - Dimensionality reduction
  - Signal denoising
  - Compression
  - ...

Google  
AdWords

Gmail<sup>™</sup>  
by Google



YouTube


# Column Subset Selection (CSS)

- Columns often have important meaning
- **CSS**: Low-rank matrix approximation in column space of  $A$

$$\arg \min_{S \subset [n], |S|=k} \|A - \Pi_{A[S]} A\|_F^2$$

The diagram illustrates the Column Subset Selection (CSS) process. It shows a matrix  $A$  of size  $m \times n$  (with  $n=10$  columns and  $m$  rows) being approximated by the product of a matrix  $A[S]$  of size  $m \times k$  (with  $k=3$  columns) and a projection matrix  $\Pi_{A[S]}$  of size  $k \times n$  (with  $n=10$  columns). The matrix  $A$  is represented by a grid of 10 columns, with 3 columns highlighted in blue. The matrix  $A[S]$  is represented by a grid of 3 columns, all highlighted in blue. The projection matrix  $\Pi_{A[S]}$  is represented by a grid of 10 columns, all highlighted in blue. The approximation is indicated by a tilde symbol ( $\approx$ ) between  $A$  and the product  $A[S] \Pi_{A[S]}$ .

# DISTGREEDY: GCSS( $A, B, k$ ) with $L$ machines




Machine 1

Machine 2

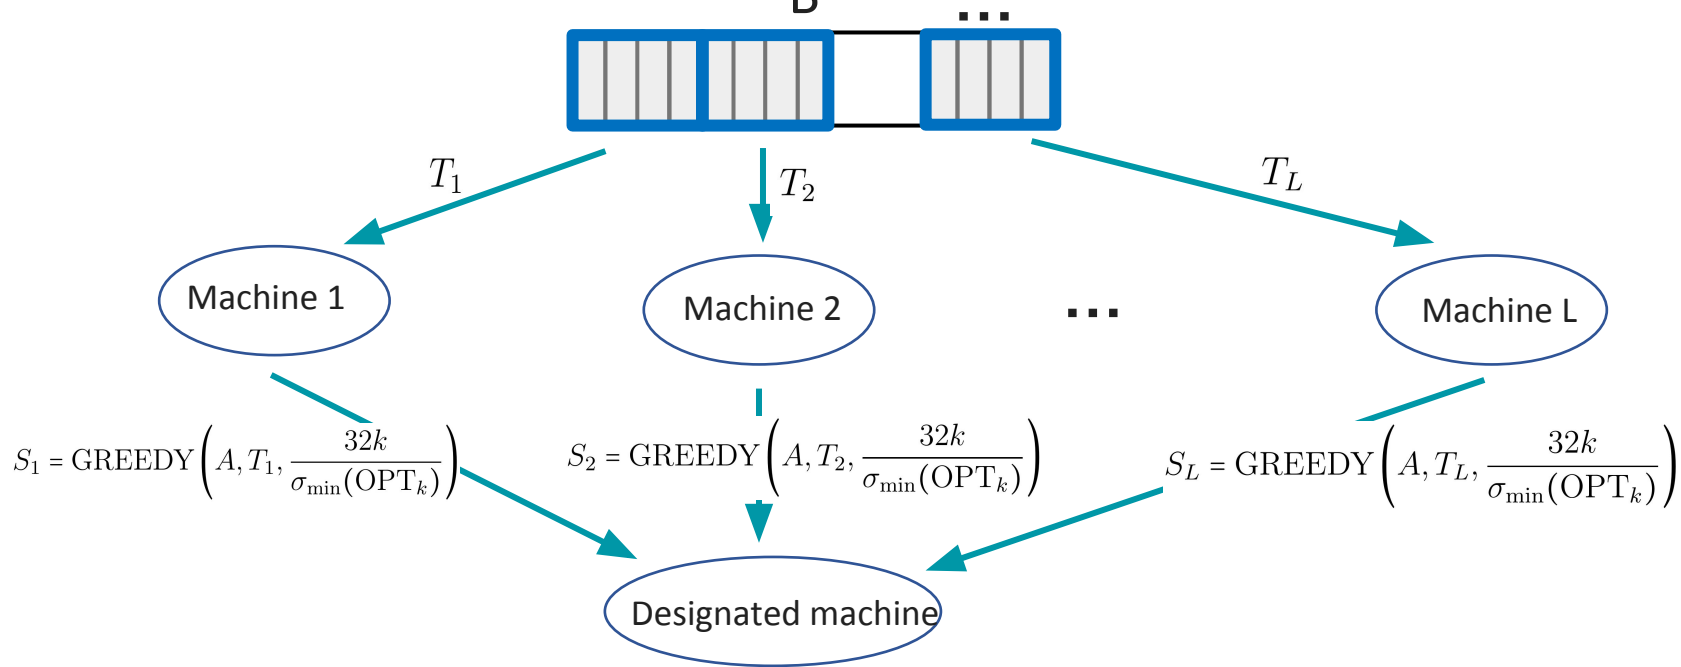
$\dots$

Machine  $L$

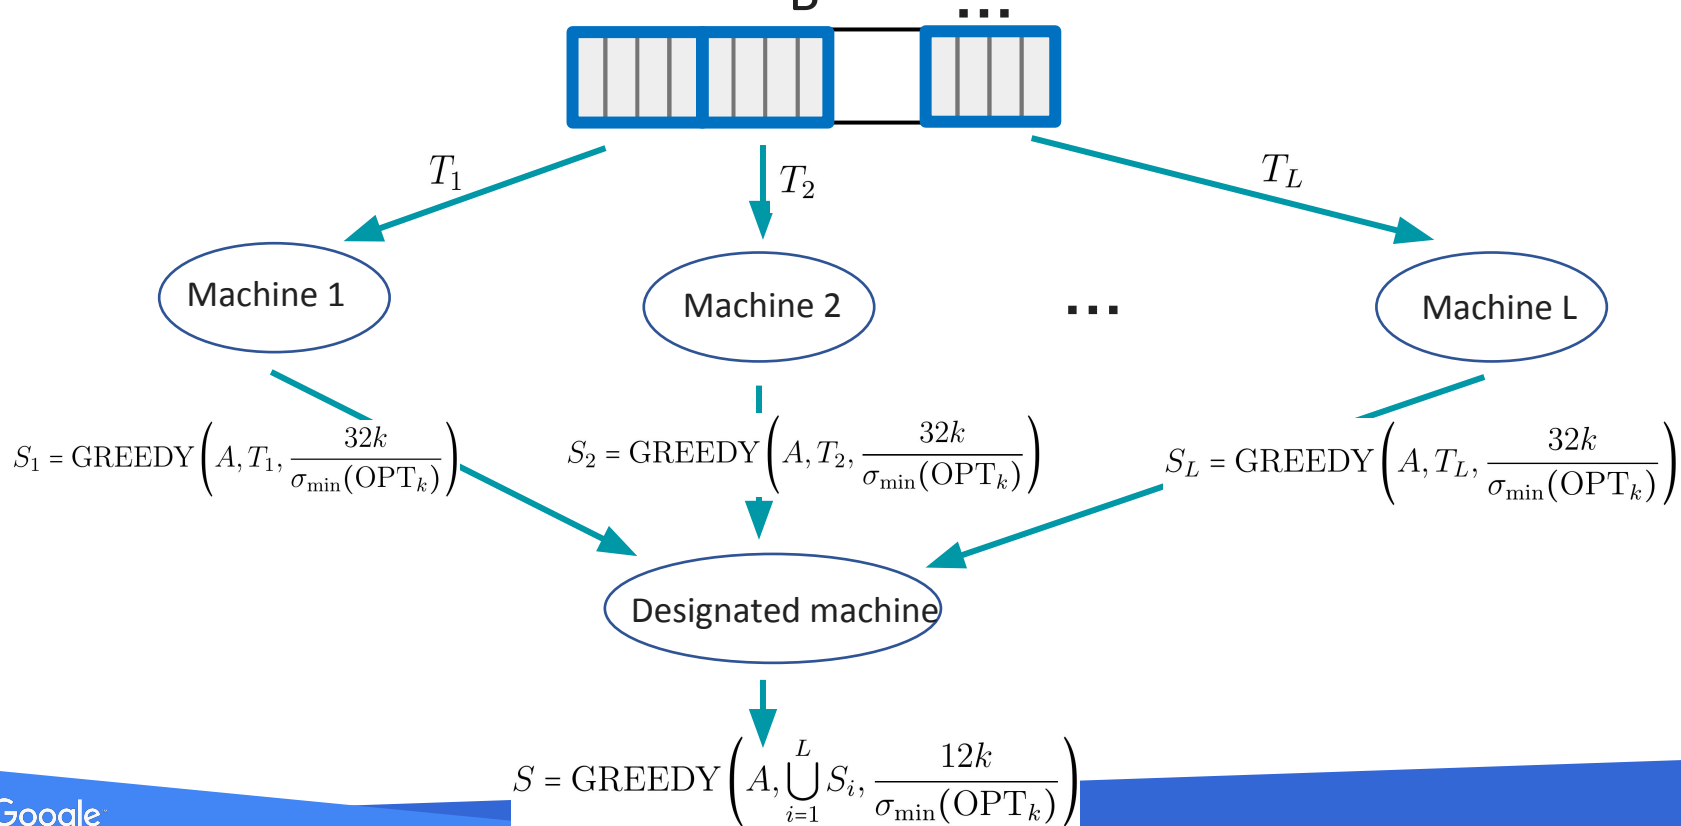
# DISTGREEDY: GCSS( $A, B, k$ ) with $L$ machines



# DISTGREEDY: GCSS(A,B,k) with L machines



# DISTGREEDY: GCSS(A,B,k) with L machines



# DISTGREEDY for column subset selection

**1 round result:** DISTGREEDY with  $r = O\left(\frac{k}{\sigma_{\min}(OPT)}\right)$  gives objective value  $\Omega\left(\frac{f(OPT_k)}{\kappa(OPT_k)}\right)$

Condition number  $\frac{\sigma_{\max}(OPT_k)}{\sigma_{\min}(OPT_k)}$

**Multi-round result:**  $O\left(\frac{\kappa(OPT)}{\varepsilon}\right)$  rounds gives objective value  $\Omega\left((1 - \varepsilon)f(OPT_k)\right)$

# Empirical result for column subset selection

- Training accuracy on massive data set (news 20.binary, 15k x 100k matrix)
- Speedup over 2-phase algorithm in parentheses

<b>n</b>	<b>Rand</b>	<b>2-Phase</b>	<b>DISTGREEDY</b>	<b>PCA</b>
500	54.9	81.8 (1.0)	80.2 (72.3)	85.8 (1.3)
1000	59.2	84.4 (1.0)	82.9 (16.4)	88.6 (1.4)
2500	67.6	87.9 (1.0)	85.5 (2.4)	90.6 (1.7)

- **Interesting experiment:** What if we partition more carefully and not randomly?
  - **Recent observation:** If we treat each machine separately, it does not help much! Random partitioning is good even compared with more careful partitioning.



# Coverage Problems

**Problems**: Given a set system ( $n$  sets and  $m$  elements),

1. “**K-coverage**”: pick  $k$  sets to max. size of union
2. “**set cover**”: cover *all* elements with least number of sets
3. “**set cover with outliers**”: cover  $(1-\lambda)m$  elements with least number of sets

# Coverage Problems

**Problems**: Given a set system ( $n$  sets and  $m$  elements),

1. “**K-coverage**”: pick  $k$  sets to max. size of union
2. “**set cover**”: cover *all* elements with least number of sets
3. “**set cover with outliers**”: cover  $(1-\lambda)m$  elements with least number of sets

Greedy Algorithm: Pick a subset with the maximum marginal coverage,

# Coverage Problems

**Problems:** Given a set system ( $n$  sets and  $m$  elements),

1. “**K-coverage**”: pick  $k$  sets to max. size of union
2. “**set cover**”: cover *all* elements with least number of sets
3. “**set cover with outliers**”: cover  $(1-\lambda)m$  elements with least number of sets

Greedy Algorithm: Pick a subset with the maximum marginal coverage,

- $1-1/e$ -approx. To  $k$ -coverage,  $\log n$ -approximation for set cover...
- Goal: Achieve good fast approximation with minimum memory footprint
  - Streaming: elements arrive one by one, not sets
  - Distributed: linear communication and memory independent of the size of ground set

# Submodular Maximization vs. Maximum Coverage

**Coverage function is a special case of submodular function:**

**$f(R)$  = cardinality of union of family  $R$  of subsets**

$$f(R) = \left| \bigcup_{S \in R} S \right|$$

# Submodular Maximization vs. Maximum Coverage

Coverage function is a special case of submodular maximization:

$f(R)$  = cardinality of union of family  $R$  of subsets

$$f(R) = \left| \bigcup_{S \in R} S \right|$$

So problem solved?

[MirrokniZadimoghaddam STOC'15]: Randomized composable core-sets work

[Mirzasoleiman et al NIPS'14]: This method works well in Practice!

# Submodular Maximization vs. Maximum Coverage

Coverage function is a special case of submodular maximization:

$f(R)$  = cardinality of union of family  $R$  of subsets

$$f(R) = \left| \bigcup_{S \in R} S \right|$$

So problem solved?

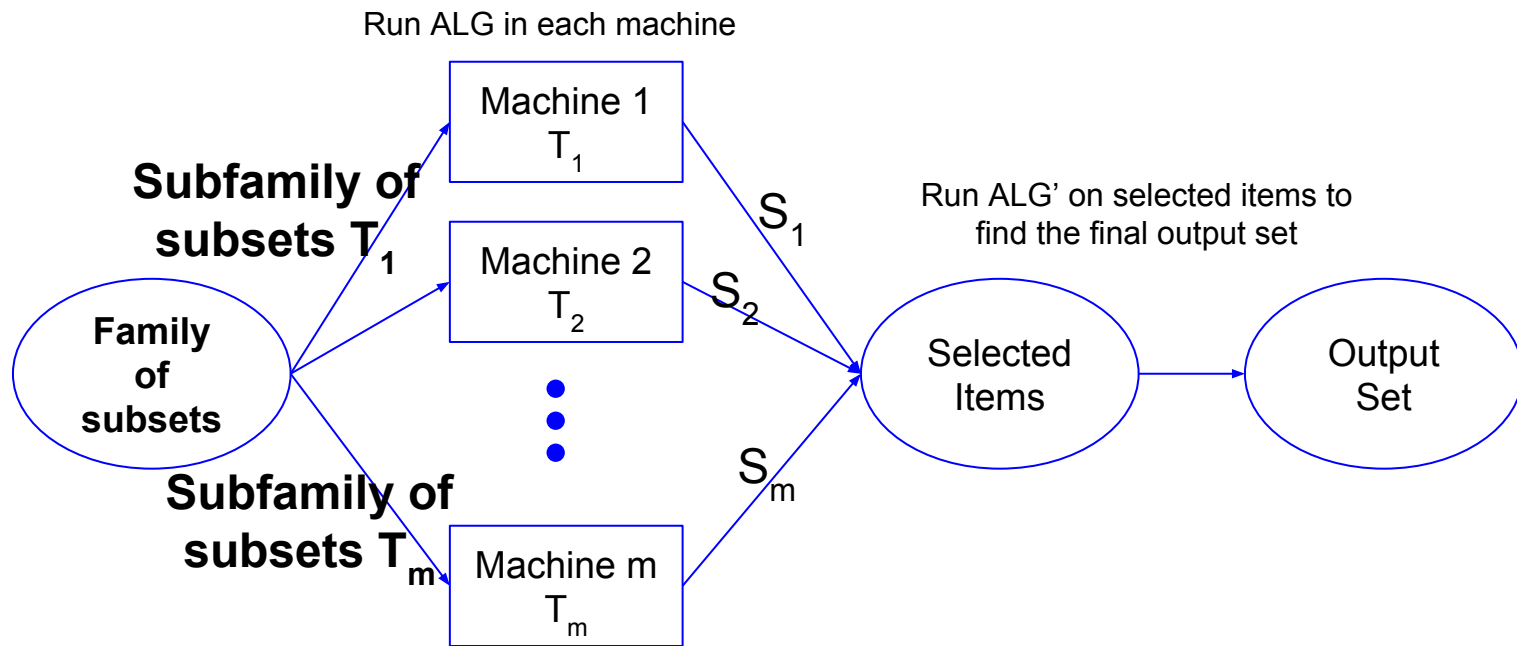
[MirrokniZadimoghaddam STOC'15]: Randomized composable core-sets work

[Mirzasoleiman et al NIPS'14]: This method works well in Practice!

**No. This solution has several issues for coverage problems:**

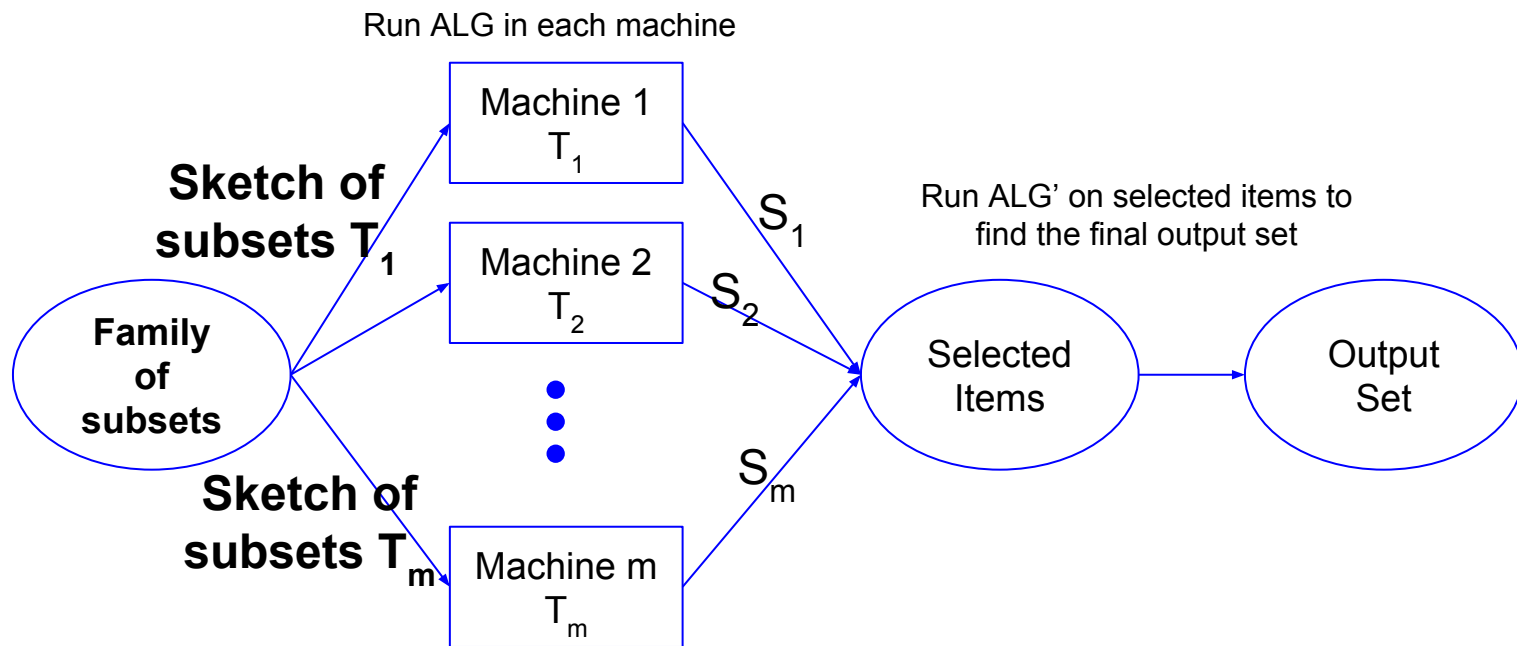
- It requires **expensive oracle access** to computing cardinality of union!
- *Distributed Computation: Send whole “sets” around ?*
- *Streaming: Handles set arrival model, does not handle “element” arrival model!*

# Why can't we apply core-sets for submodular functions?



What if the subsets are large? Can we send a sketch of them?

Idea: Send a sketch for each set (e.g., sample of elements)



Question: Does any approximation-preserving sketch work?



# Approximation-preserving sketching is not sufficient.

Idea: Use sketching to define **a  $(1 \pm \epsilon)$ -approx oracle to cardinality of union function?**

**[BateniEsfandiarMirrokni'16]:**

- Thm 1: A  $(1 \pm \epsilon)$ -approx sketch of coverage function **May NOT Help**
  - Given an  $(1 \pm \epsilon)$ -approx oracle to coverage function, we get  $n^{0.49}$  approximation

# Approximation-preserving sketching is not sufficient.

Idea: Use sketching to define **a  $(1 \pm \epsilon)$ -approx oracle to cardinality of union function?**

**[BateniEsfandiariMirrokni'16]:**

- Thm 1: A  $(1 \pm \epsilon)$ -approx sketch of coverage function **May NOT Help**
  - Given an  $(1 \pm \epsilon)$ -approx oracle to coverage function, we get  $n^{0.49}$  approximation
- Thm 2: With some tricks, MinHash-based sketch + proper sampling **WORKS**
  - Sample elements not sets (different from previous coresets idea)
  - Correlation between samples (MinHash)
  - Cap degrees of elements in the sketch (reduces memory footprint)


# Bipartite Graph Formulation for Coverage Problems

Bipartite graph  $G(U, V, E)$

$U$ : sets  
 $V$ : elements  
 $E$ : membership

*Set cover problem*: Pick minimum number of sets that cover all elements.

*Set cover with outliers problem*: Pick minimum number of sets that cover a  $1 - \lambda$  fraction of elements.



*Maximum coverage problem*: Pick  $k$  sets that cover maximum number of elements.

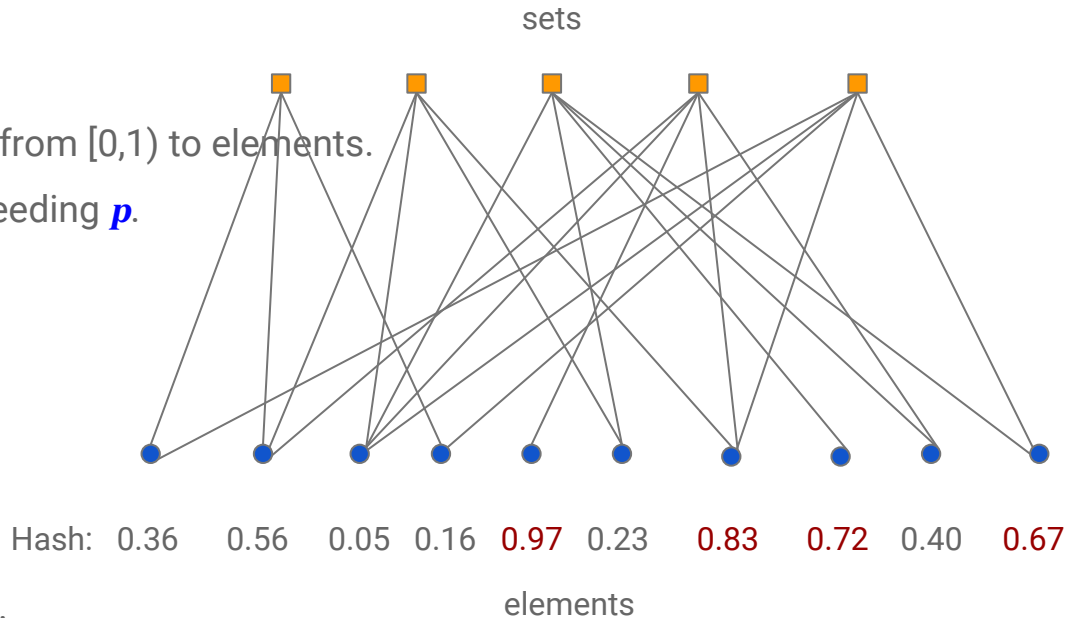
# Sketching Technique

## Construction

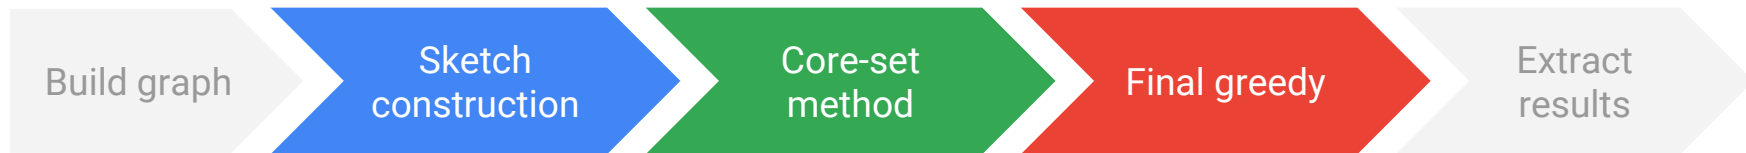
- Dependent sampling: Assign hash values from  $[0,1)$  to elements.
- Remove any element with hash value exceeding  $p$ .
- Arbitrarily remove edges to have max-degree  $\Delta$  for elements.

## Parameters

- 1)  $\Delta$  is easy to compute.
- 2)  $p$  can be found via a round of MapReduce.



# Approach



Sketch: sparse subgraph with sufficient information

For instance with many sets, parallelize using core sets.

Any single-machine greedy algorithm

## *Proof ingredients:*

1. Parameters are chosen to produce small sketch (indep. of size of ground set):  $O(\#sets)$ 
  - Challenge: how to choose parameters in distributed or streaming models
2. Any  $\alpha$ -approximation on the sketch is an  $\alpha + \epsilon$  approximation for original instance

# Summary of Results for Coverage Functions



- Special case of submodular maximization
- Problems are **NP-hard** and **APX-hard**
- Greedy algorithm gives best guarantees



Good implementations (linear-time)

- Lazy greedy algorithm
- Lazier-than-lazy algorithm

**Problem:** Graph should be stored in RAM



**Our algorithm:**

- Memory  $O(\#\text{sets})$
- Linear-time
- Optimal approximation guarantees
- MapReduce, streaming, etc.

## GREEDY

- 1) Start with empty solution
- 2) Until “done,”
  - (a) find set with best marginal coverage, and
  - (b) add it to tentative solution.

# Bounds for distributed coverage problems

From [BEM'16]: 1) Space indep. of size of sets or ground set, 2) Optimal Approximation Factor, 3) Communication linear in #sets (indep. of their size), 4) small #rounds

Previous work: [39]=[CKT'11], [42]=[MZ'15], [19]=[BENW'16], [43]=[MBKK'16]

Problem	Credit	# rounds	Approximation	Load per machine	Comment
$k$ -cover	[39]	$O(\frac{1}{\varepsilon\delta} \log m)$	$1 - \frac{1}{e} - \varepsilon$	$O(mkn^\delta)$	submodular functions
$k$ -cover	[42]	2	0.54	$\max(mk^2, mn/k)$	submodular functions
$k$ -cover	[19]	$\frac{1}{\varepsilon}$	$1 - \frac{1}{e} - \varepsilon$	$\frac{\max(mk^2, mn/k)}{\varepsilon}$	submodular functions
$k$ -cover	Here	3	$1 - \frac{1}{e} - \varepsilon$	$\tilde{O}(n + m)$	-
Set cover w outliers	Here	3	$(1 + \varepsilon) \log \frac{1}{\lambda}$	$\tilde{O}(n + m)$	-
Set cover	[43]	$\log(nm)$	$(1 + \varepsilon) \log n$	$\Omega(mn^{1-\varepsilon})$	Submodular cover
Set cover	Here	$r$	$(1 + \varepsilon) \log n$	$\tilde{O}(nm^{O(\frac{1}{r})} + m)$	-



# Bounds for streaming coverage problems

From [BEM'16]: 1) Space indep. of size of ground set, 2) Optimal Approximation Factor, 3) "Edge" vs "set" arrival

Previous work:[14]=[CW'15], [22]=[DIMV'14], [24]=[ER'14], [31]=[IMV'15], [49]=[SG'09]

Problem	Credit	# passes	Approximation	Space	Arrival
$k$ -cover	[49]	1	$1/4$	$\tilde{O}(m)$	set
$k$ -cover	Here	1	$1 - 1/e - \varepsilon$	$\tilde{O}(n)$	edge
Set cover w outliers	[24, 14]	$p$	$O(\min(n^{\frac{1}{p+1}}, e^{-\frac{1}{p}}))$	$\tilde{O}(m)$	set
Set cover w outliers	Here	1	$(1 + \varepsilon) \log \frac{1}{\lambda}$	$\tilde{O}_\lambda(n)$	edge
Set cover	[14, 49]	$p$	$(p + 1)n^{\frac{1}{p+1}}$	$\tilde{O}(m)$	set
Set cover	[22]	$4^k$	$4^k \log n$	$\tilde{O}(nm^{\frac{1}{k}})$	set
Set cover <sup>1</sup>	[31]	$p$	$O(p \log n)$	$\tilde{O}(nm^{O(\frac{1}{p})})$	set
Set cover	Here	$p$	$(1 + \varepsilon) \log n$	$\tilde{O}(nm^{O(\frac{1}{p})} + m)$	edge

# Empirical Study

## Public datasets

- Social networks
- Bags of words
- Contribution graphs
- Planted instances

Name	Type	$ S $	$ E $	$ E $
livejournal-3	dominating set	3,997,962	3,997,962	72,803,204,325
livejournal-2	dominating set	3,997,962	3,997,962	3,377,182,611
dblp-3	dominating set	317,080	317,080	333,505,724
dblp-2	dominating set	317,080	317,080	27,437,914
gutenberg	bag of words	41,716	99,949,091	1,068,977,156
s-gutenberg	bag of words	925	10,620,424	27,337,479
reuters	bag of words	199,328	138,922	15,334,605
planted-A	planted	10,100	10,000	1,220,000
planted-B	planted	100,100	1,000,000	1,201,100,000
planted-C	planted	100,500	10,000,000	2,410,100,000
planted-D	planted	101,000	10,000,000	1,210,100,000
wiki-main	contribution graph	2,953,425	10,619,081	75,151,304
wiki-talk	contribution graph	1,736,343	1,017,617	7,299,920

Instance	Footprint	Quality
wiki-main	0.06%	94.4%
wiki-main	2.4%	99.5%
wiki-main	7.7%	99.9%
wiki-talk	1.5%	99.2%
planted-A	8.2%	96%

Instance	Footprint	Quality
reuters	10%	96%
dblp-2	1.7%	92%
dblp-2	3.1%	96%
reuters	1.2%	87%
reuters	3.6%	92%

– Very small sketches (0.01–5%) suffice for obtaining good approximations (95+%).

– Without core sets, can handle in <1h XXXB edges or elements.

# Feature Selection (ongoing)

**Goal**: Pick k “representative” features

Based on composable core sets

k	Random clusters	Best cluster method	Set cover (pairs)
500	0.8538	0.851	0.862
1000	0.8864	0.8912	0.8936
2500	0.9236	0.9234	0.9118

features

entities

+	+	+	+	+	+	+	+	+	+
+	+	+	+	-	-	-	-	-	-
+	-	-	-	+	+	+	-	-	-
-	+	-	-	+	-	-	+	+	-
-	-	+	-	-	+	-	+	-	+
-	-	-	+	-	-	+	-	+	+

- 1) Pick features that cover all entities
- 2) Pick features that cover many pairs (or triples, etc.) of entities

# Summary: Distributed Algorithms for Five Problems

Define on a metric space & composable core-sets apply.

1. Diversity Maximization,
  - PODS'14 by IndykMahdianMahabadiM.
  - for Feature Selection in AAAI'17 by AbbasiGhadiriMirrokniZadimoghaddam
2. Capacitated  $\ell_p$  Clustering, NIPS'14 by BateniBhaskaraLattanziM.

Beyond Metric Spaces. Only *Randomized* partitioning apply.

3. Submodular Maximization, STOC'15 by M. Zadimoghaddam
4. Feature Selection (Column Subset Selection), ICML'16 by Alschulter et al.

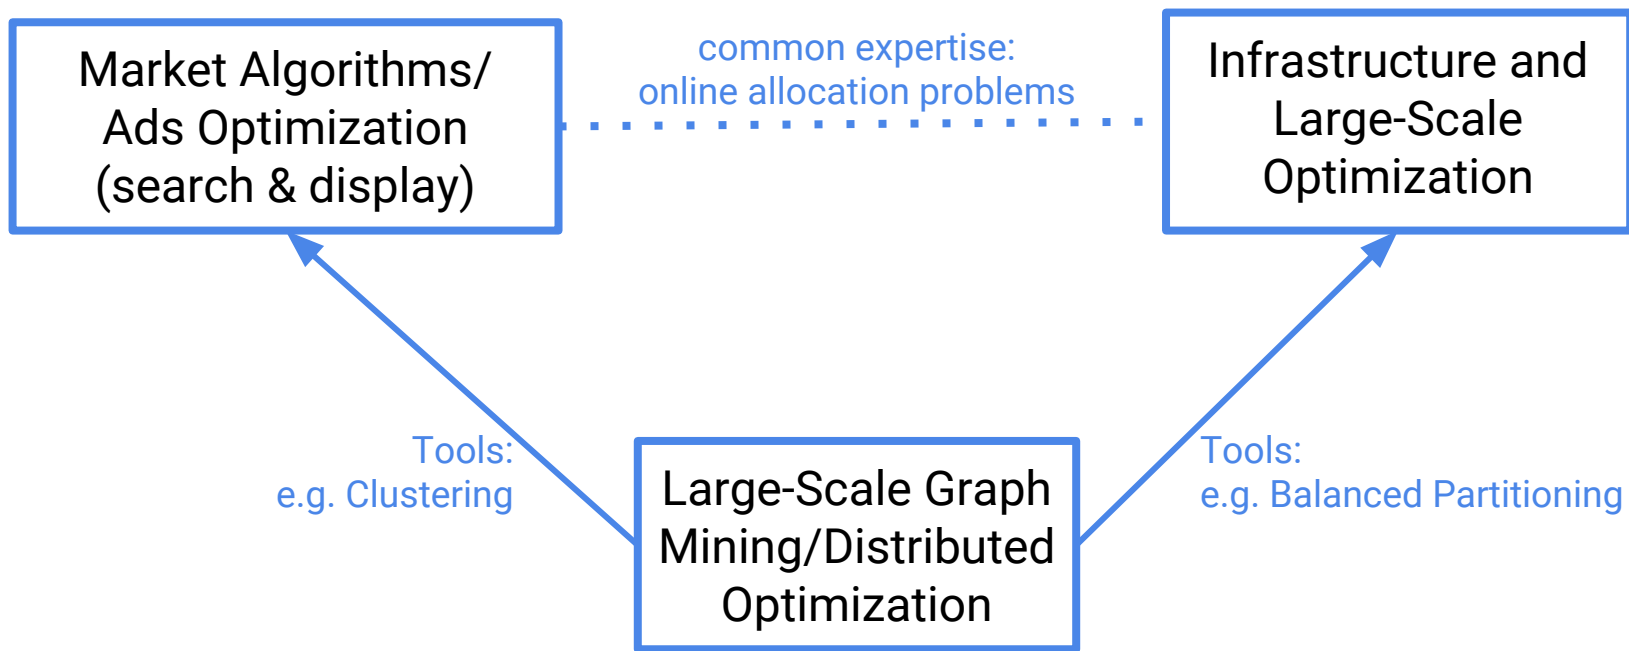
Needs adaptive sampling/sketching techniques

5. Coverage Problems: by BateniEsfandiariM




# Our team: Google NYC Algorithms Research Team

Recently released external team website: [research.google.com/teams/nycalg/](https://research.google.com/teams/nycalg/)



# THANK YOU



mirrokni@google.com



# Local Search for Diversity Maximization [KDD'13]

- Used for sum of pairwise distances
- Algorithm [Abbasi, Mirrokni, Thakur]
  - Initialize  $S$  with an arbitrary set of  $k$  points which contains the two farthest points
  - While there exists a swap that improves diversity by a factor of  $\left(1 + \frac{\epsilon}{n}\right)$ 
    - » Perform the swap
- For Remote-Clique
  - Number of rounds:  $\log_{\left\{1 + \frac{\epsilon}{n}\right\}} k^2 = O\left(\frac{n}{\epsilon} \log k\right)$
  - Approximation factor is constant.

