

Parallel Peeling Algorithms

Justin Thaler, Yahoo Labs

Joint Work with:

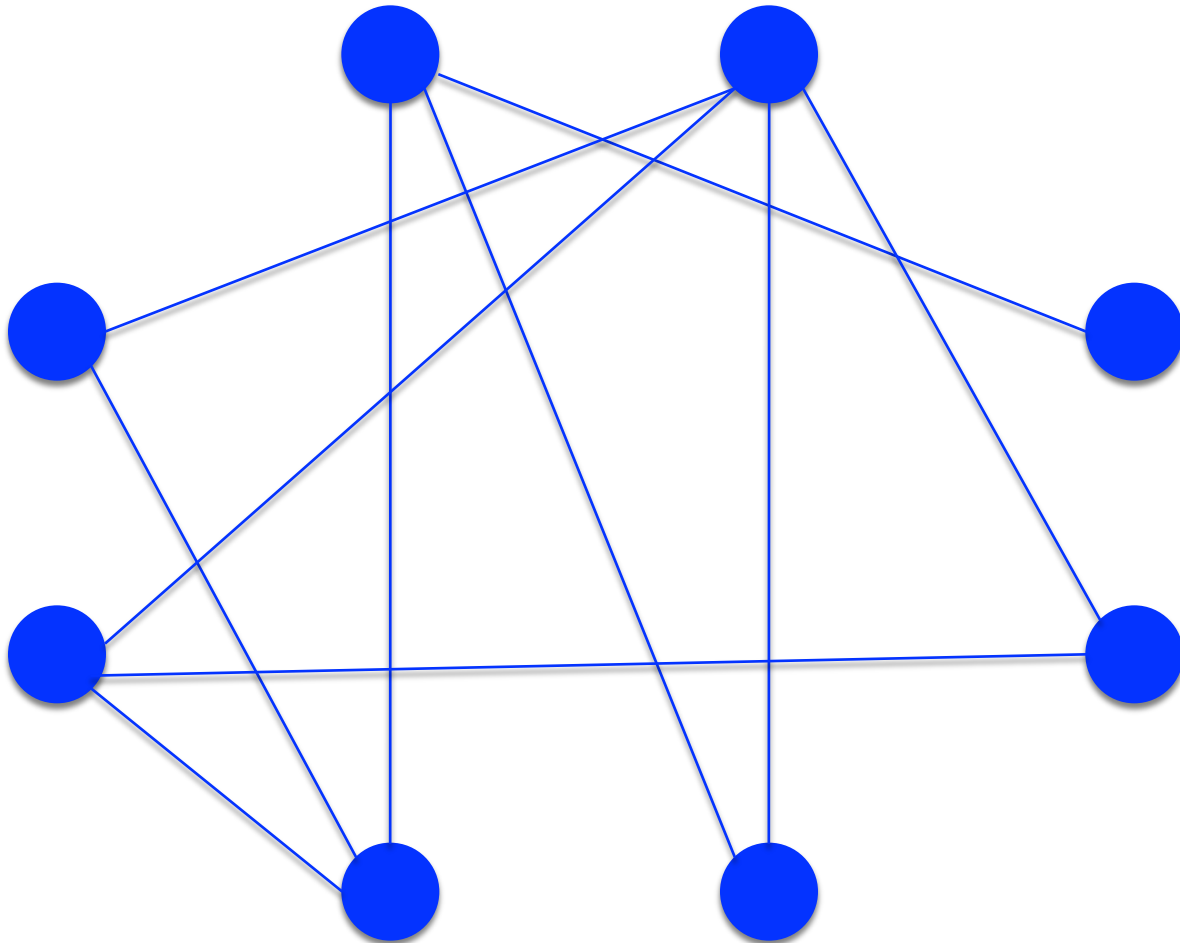
Michael Mitzenmacher, Harvard University

Jiayang Jiang

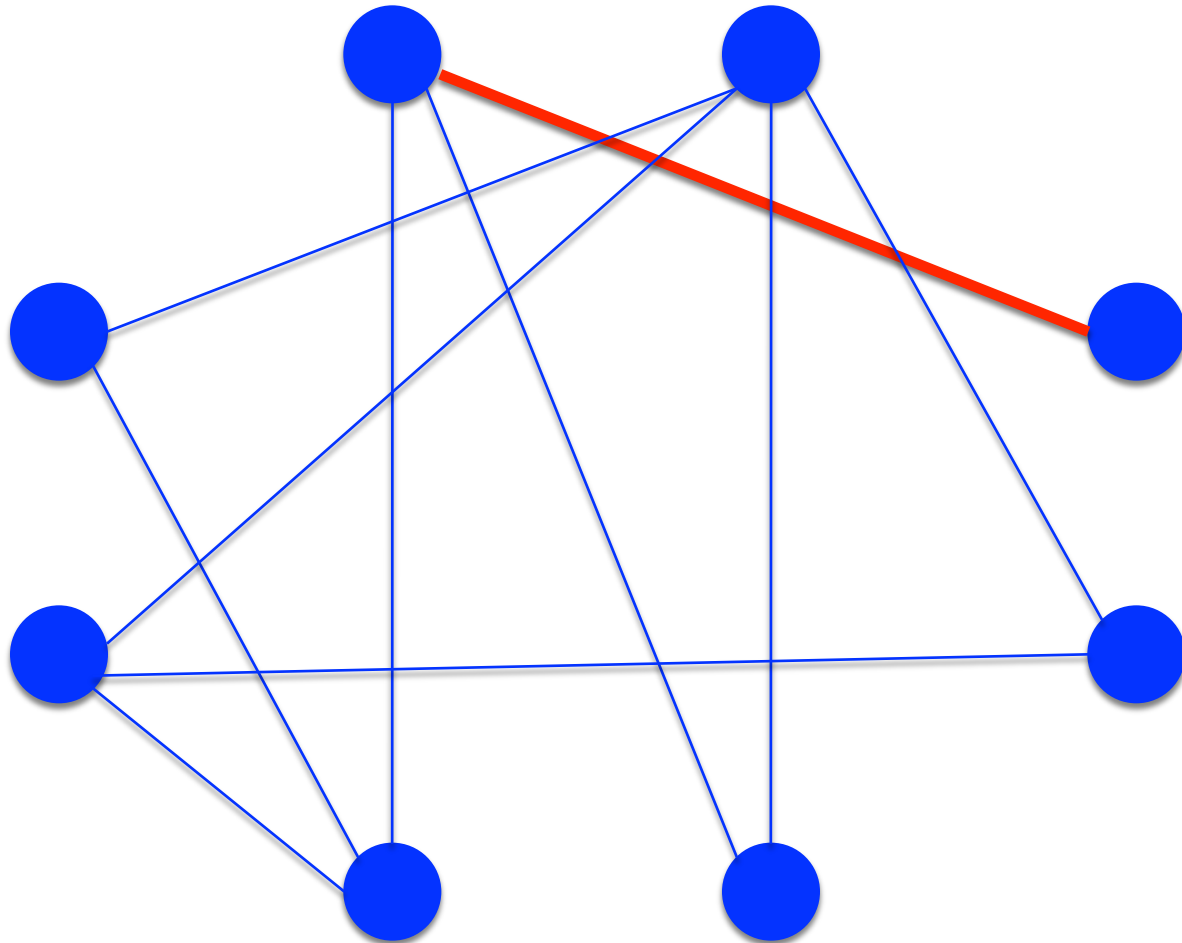
The Peeling Paradigm

- Many important algorithms for a wide variety of problems can be modeled in the same way.
- Start with a (random) hypergraph G .
 - While there exists a node v of degree less than k :
 - Remove v and all incident edges.
- The remaining graph is called the **k -core** of G .
 - $k=2$ in most applications.
- Typically, the algorithm “succeeds” if the the k -core is empty.
 - To ensure “success”, data structure should be designed large enough so that the k -core of G is empty w.h.p.
- Typically yields simple, greedy algorithms running in linear time.

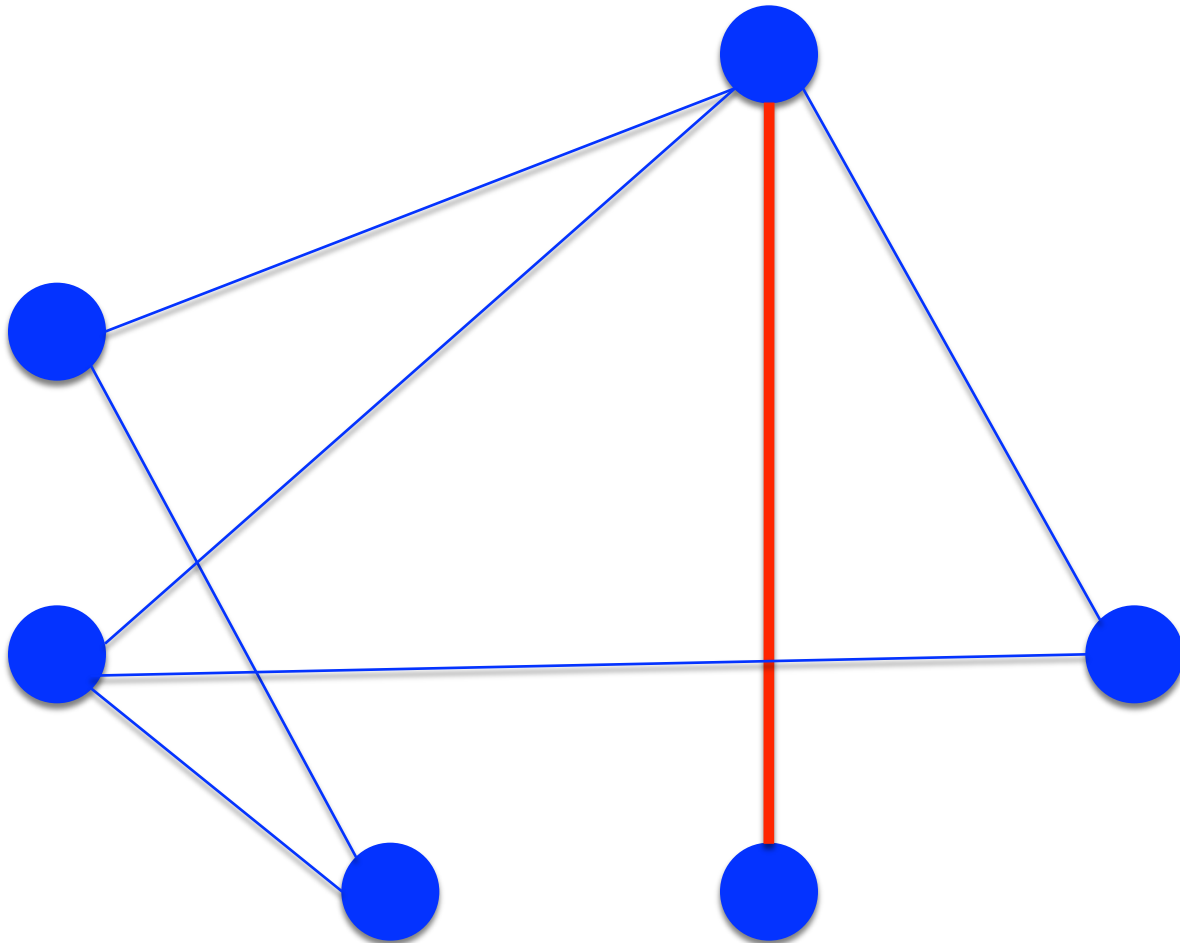
The peeling process when $k=2$



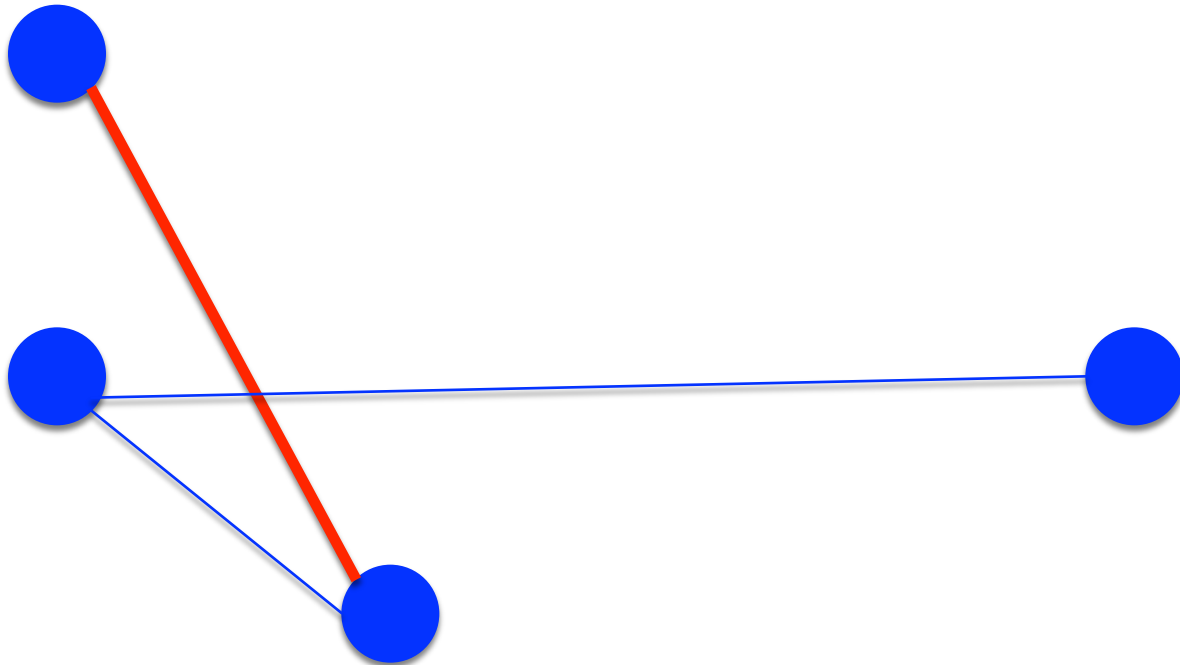
The peeling process when $k=2$



The peeling process when $k=2$



The peeling process when $k=2$



The peeling process when $k=2$



Example Algorithms

Example 1: Sparse Recovery Algorithms

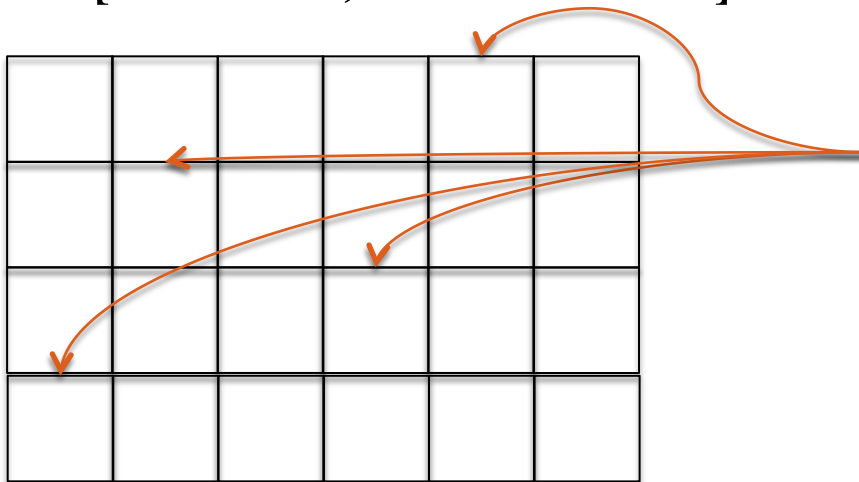
- Consider data streams that insert and delete a lot of items.
 - Flows through a router, people entering/leaving a building.
- Sparse Recovery problem: list all items with non-zero frequency.
- Want listing not at all times, but at “reasonable” or “off-peak” times, when working set size is bounded.
 - If we do M insertions, then $M-N$ deletions, and want a list at the end, we need to list N items.
- Data structure size should be proportional to N , not to M !
 - Proportional to size you want to be able to list, not number of items your system has to handle.
- Central primitive used in more complicated streaming algorithms.
 - E.g. L_0 sampling, which is in turn used to solve problems on dynamic graph streams (see previous talk).

Example 1: Sparse Recovery Algorithms

- For simplicity, assume that when listing occurs, no item has frequency more than 1.

Example 1: Sparse Recovery Algorithms

- Sparse Recovery Algorithm: Invertible Bloom Lookup Tables (IBLTs)
[Goodrich, Mitzenmacher]



Each stream item hashed to r cells
(using r different hash functions)

Count
KeySum

Insert(x): For each of the j cells that x is hashed to:

Add key to KeySum

Increment Count

Delete(x): For each of the j cells x is hashed to:

Subtract key from keysum

Decrement Count

Listing Algorithm: Peeling

- Call a cell “pure” if its count equals 1.
- While there exists a pure cell:
 - Output $x = \text{keySum}$ of the cell.
 - Call $\text{Delete}(x)$ on the IBLT.

Listing Algorithm: Peeling

- Call a cell “pure” if its count equals 1.
- While there exists a pure cell:
 - Output $x = \text{keySum}$ of the cell.
 - Call $\text{Delete}(x)$ on the IBLT.
- To handle frequencies that are larger than 1, add a checksum field to each cell (details omitted).

Listing Algorithm: Peeling

- Call a cell “pure” if its count equals 1.
- While there exists a pure cell:
 - Output $x = \text{keySum}$ of the cell.
 - Call $\text{Delete}(x)$ on the IBLT.
- To handle frequencies that are larger than 1, add a checksum field to each cell (details omitted).
- Listing \longleftrightarrow peeling to 2-core on the hypergraph G where:
 - Cells \longleftrightarrow vertices of G .
 - Items in IBLT \longleftrightarrow hyperedges of G .
 - G is r -uniform (each edge has r vertices, one for each cell the item is hashed to).

How Many Cells Does an IBLT Need to Guarantee Successful Listing?

- Consider a random r -uniform hypergraph G with n nodes and $m = c \cdot n$ edges.
 - i.e., each edge has r vertices, chosen uniformly at random from $[n]$ without repetition.
- Known fact: Appearance of a non-empty k -core obeys a sharp threshold.
 - For some constant $c_{k,r}$, when $m < c_{k,r}n$, the k -core is empty with probability $1 - o(1)$.
 - When $m > c_{k,r}n$, the k -core of G is non-empty with probability $1 - o(1)$.
 - Implication: to successfully list a set of size M with probability $1 - o(1)$, the IBLT needs roughly $M / c_{k,r}$ cells.
 - E.g. $c_{2,3} \approx 0.818$, $c_{2,4} \approx 0.772$, $c_{3,3} \approx 1.553$.

How Many Cells Does an IBLT Need to Guarantee Successful Listing?

- Consider a random r -uniform hypergraph G with n nodes and $m = c \cdot n$ edges.
 - i.e., each edge has r vertices, chosen uniformly at random from $[n]$ without repetition.
- Known fact: Appearance of a non-empty k -core obeys a sharp threshold.
 - For some constant $c_{k,r}$, when $m < c_{k,r}n$, the k -core is empty with probability $1 - o(1)$.
 - When $m > c_{k,r}n$, the k -core of G is non-empty with probability $1 - o(1)$.
 - Implication: to successfully list a set of size M with probability $1 - o(1)$, the IBLT needs roughly $M / c_{k,r}$ cells.
 - E.g. $c_{2,3} \approx 0.818$, $c_{2,4} \approx 0.772$, $c_{3,3} \approx 1.553$.
 - In general:

$$c_{k,r}^* = \min_{x>0} \frac{x}{r(1 - e^{-x} \sum_{j=0}^{k-2} \frac{x^j}{j!})^{r-1}}.$$

Other Examples of Peeling Algorithms

- Low-Density Parity Check Codes for Erasure Channel.
 - [Luby, Mitzenmacher, Shokrollah, Spielman]
- Biff codes (directly use IBLTs).
 - [Mitzenmacher and Varghese]
- k -wise independent hash families with $O(1)$ evaluation time.
 - [Siegel]
- Sparse FFT algorithms.
 - [Hassanieh et al.]
- Cuckoo hashing.
 - [Pagh and Rodler]
- Pure literal rule for computing satisfying assignments of random CNFs.
 - [Franco] [Mitzenmacher] [Molloy] [many others].

Parallel Peeling Algorithms

Our Goal: Parallelize These Peeling Algorithms

- Recall: the aforementioned algorithms are equivalent to peeling a random hypergraph G to its k -core.
- There is a brain dead way to parallelize the peeling process.
 - For each node v in parallel:
 - Check if v has degree less than k .
 - If so, remove v and its incident hyperedges.
- Key question: how many rounds of peeling are required to find the k -core?
- Algorithm is simple, analysis is tricky.

Main Result

- Two behaviors:
 - Parallel peeling completes in $O(\log \log n)$ rounds if the edge density c is “below the threshold” $c_{k,r}$.
 - Parallel peeling requires $\Omega(\log n)$ rounds if the edge density c is “above the threshold” $c_{k,r}$.
- This is great!
 - Most peeling uses the goal is to be *below the threshold*.
 - So “nature” is helping us by making parallelization fast.
 - Implies $\text{poly}(\log \log n)$ time, $O(n \text{ poly}(\log \log n))$ work, parallel algorithms for listing elements in an IBLT, decoding LDPC codes, etc.

Precise Upper Bound

Theorem 1. *Let $k, r \geq 2$ with $k + r \geq 5$, and let c be a constant. With probability $1 - o(1)$, the parallel peeling process for the k -core in a random hypergraph $G_{n,cn}^r$ with edge density c and r -ary edges terminates after $\frac{1}{\log((k-1)(r-1))} \log \log n + O(1)$ rounds when $c < c_{k,r}^*$.*

Theorem 2. *Let $k, r \geq 2$ with $k + r \geq 5$, and let c be a constant. With probability $1 - o(1)$, the parallel peeling process for the k -core in a random hypergraph $G_{n,cn}^r$ with edge density c and r -ary edges requires $\frac{1}{\log((k-1)(r-1))} \log \log n - O(1)$ rounds to terminate when $c < c_{k,r}^*$.*

Summary: The right factor in front of the $\log \log n$ is $1 / (\log(k-1)(r-1))$
(tight up to an additive constant).

Lower Bound

Theorem 3. *Let $r \geq 3$ and $k \geq 2$. With probability $1 - o(1)$, the peeling process for the k -core in $G_{n,cn}^r$ terminates after $\Omega(\log n)$ rounds when $c > c_{k,r}^*$.*

Summary: $\Omega(\log n)$ lower bound matches an earlier $O(\log n)$ upper bound due to [Achlioptas and Molloy, 2013].

Proof Sketch for Upper Bound

- Let λ_i denote the probability a given vertex v survives i rounds of peeling.
- Claim: $\lambda_{i+1} \leq (C\lambda_i)^{(k-1)(r-1)}$ for some constant C .
 - Suggests $\lambda_i \ll 1/n$ after about $1 / ((k-1)(r-1)) * \log \log n$ rounds.
 - A related argument shows that $\lambda_i \leq 1/(2C)$ after $O(1)$ rounds, and after that point the claim implies that λ_i falls doubly-exponentially quickly.

Proof Sketch for Upper Bound

- Let λ_i denote the probability a given vertex v survives i rounds of peeling.
- Claim: $\lambda_{i+1} \leq (C\lambda_i)^{(k-1)(r-1)}$ for some constant C .
- **Very** crude sketch of the Claim's plausibility:
 - Node v survives round $i+1$ only if it has (at least) k incident edges $e_1 \dots e_k$ that survive round i .
 - Fix a k -tuple of edges $e_1 \dots e_k$ incident to v .
 - Assume no node other than v appears in more than one of these edges.
 - Then there are $k(r-1)$ distinct nodes other than v appearing in these edges.
 - The edges all survive round i only if **all** $k(r-1)$ of these nodes survive round i .
 - Let's pretend that the survival of these nodes are independent events.
 - Then the probability all nodes survive round i is roughly $\lambda_i^{k(r-1)}$.
 - Finally, union bound over all k -tuples of edges incident to v .

Simulation Results

	$c = 0.7$		$c = 0.75$		$c = 0.8$		$c = 0.85$	
n	Failed	Rounds	Failed	Rounds	Failed	Rounds	Failed	Rounds
10000	0	12.504	0	23.352	1000	17.037	1000	10.773
20000	0	12.594	0	23.433	1000	19.028	1000	11.928
40000	0	12.791	0	23.343	1000	20.961	1000	12.992
80000	0	12.939	0	23.372	1000	22.959	1000	14.104
160000	0	12.983	0	23.421	1000	25.066	1000	15.005
320000	0	13.000	0	23.491	1000	27.089	1000	16.305
640000	0	13.000	0	23.564	1000	29.281	1000	17.334
1280000	0	13.000	0	23.716	1000	31.037	1000	18.499
2560000	0	13.000	0	23.840	1000	33.172	1000	19.570

- Results from simulations of parallel peeling process on random 4-uniform hypergraphs with n nodes and $c*n$ edges using $k = 2$.
- Averaged over 1000 trials.
- Recall that $c_{2,4} \approx 0.772$.

Refined Result: Mind the Gap

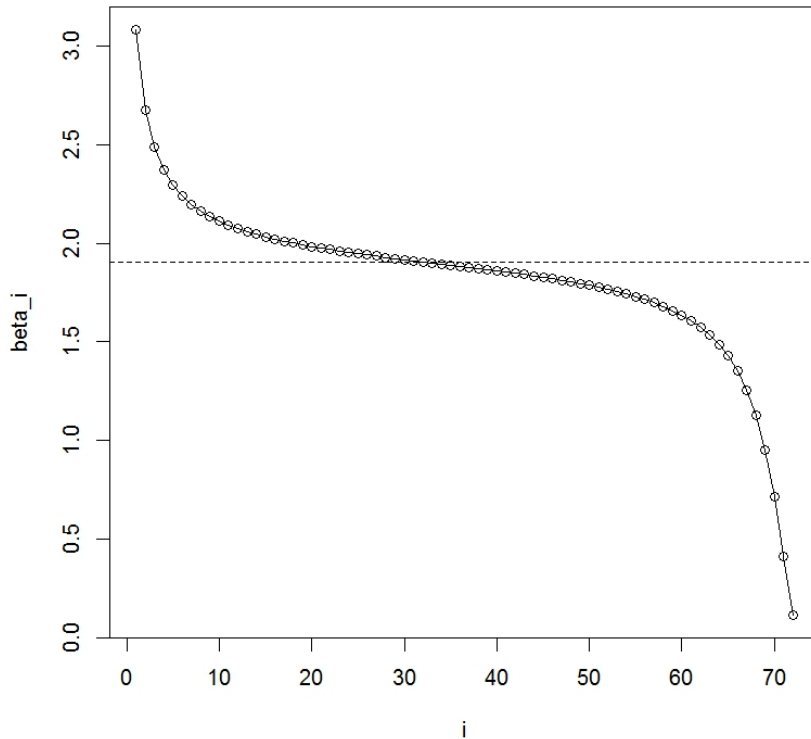
THEOREM 7.1. *Let $v = |c_{k,r}^* - c|$ for constant c with $c < c_{k,r}$. With probability $1 - o(1)$, peeling in $G_{n,cn}^r$ requires $\Theta(\sqrt{1/v}) + \frac{1}{\log((k-1)(r-1))} \log \log n$ rounds when c is below the threshold density $c_{k,r}^*$.*

Summary: below the threshold, the additive term is $\Theta(1/\sqrt{|\text{gap}|})$.

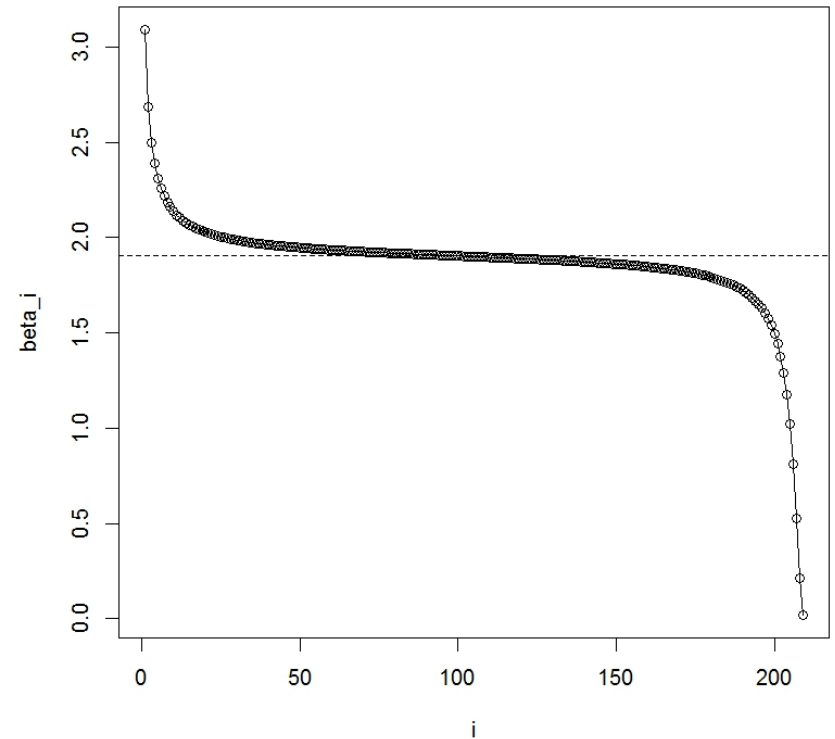
This can be more important than the $\log \log n$ term if the edge density is close to the threshold!

Refined Simulations: Mind the Gap

Plot for β_i , $r=4$, $k=2$, $c=0.77$



Plot for β_i , $r=4$, $k=2$, $c=0.772$



Plots show expected progress of the peeling process as a function of the round i , for values of the edge density c approaching the threshold value of $c_{2,4} \approx 0.772$.

Refined Analysis: Mind the Gap

- Analysis shows that peeling process falls into three “stages”.
 - First stage: the fraction of surviving nodes falls very quickly as a function of the rounds until it gets close to a certain key value x^* .
 - Second stage: $\Theta(1/\sqrt{|\text{gap}|})$ rounds are required to go from “close” to x^* to “significantly below” x^* .
 - Third stage: the analysis of the basic upper bound kicks in, and the fraction of surviving nodes falls doubly-exponentially quickly.

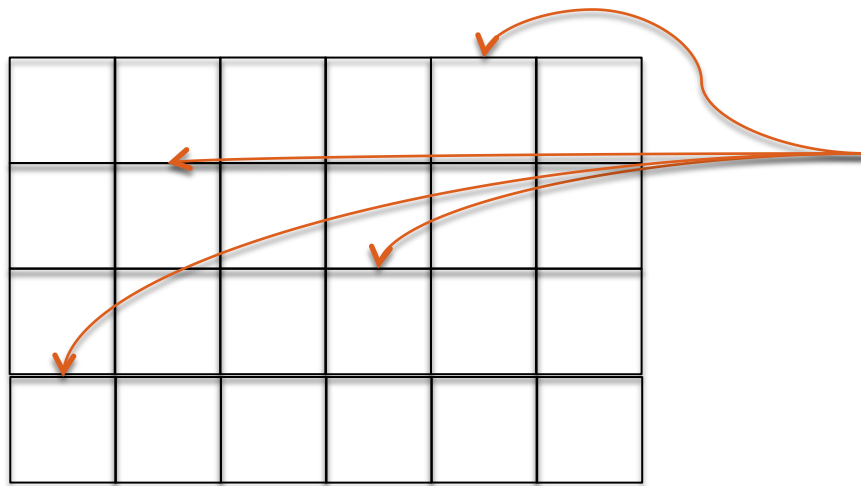
Implementation Issues

GPU Experimental Results

Table Load	No. Table Cells	% Recovered	GPU Recovery Time	Serial Recovery Time	GPU Insert Time	Serial Insert Time
0.75	16.8 million	100%	0.33 s	6.37 s	0.31 s	3.91 s
0.83	16.8 million	50.1%	0.42 s	3.64 s	0.35 s	4.34 s

Table 3: Results of our parallel and serial IBLT implementations with $r = 3$ hash functions. The table load refers to the ratio of the number of items in the IBLT to the number of cells in the IBLT.

Recall: IBLTs



Each stream item hashed to r cells
(using r different hash functions)

Count
KeySum

Insert(x): For each of the j cells that x is hashed to:

 Add key to KeySum

 Increment Count

Delete(x): For each of the j cells x is hashed to:

 Subtract key from keysum

 Decrement Count

Recall: IBLT Listing Algorithm

- Call a cell “pure” if its count equals 1.
- While there exists a pure cell:
 - Output $x = \text{keySum}$ of the cell.
 - Call $\text{Delete}(x)$ on the IBLT.

GPU Implementation

- Each cell gets a thread.
- Each cell checks if it is pure.
 - If so, identify the key it contains and remove it from other cells in the IBLT.
 - Do this by subtracting out values in other cells.
- Issue: repeated deletion.
 - Several cells might recover and try to remove the same key in the same round. So a key gets deleted more than once!

Dealing with Repeated Deletion

- To avoid this: use r subtables, such that the i th hash function only hashes into subtable i .
 - Break the listing algorithm into serial subrounds. In i th subround, recover only from the i th subtable.
 - Avoids repeated deletions, since each item will be hashed to just 1 cell in each subtable.
 - Leads to interesting variation in the analysis.
- Subrounds increase runtime, since they must happen sequentially.
 - Naively, they may blow up runtime by a factor of r .
 - But we show this does not happen.
 - Gains in one subround can help later subrounds.
 - We show runtime only blows up by a factor of about $\log_2(r-1)$.
- Analysis is similar to Vöcking's d -left scheme.
 - Fibonacci numbers show up!

Subround Result

THEOREM B.1. *Let $r \geq 3$ and $k \geq 2$. Let $\phi_{r-1} = \lim_{k \rightarrow \infty} F_{r-1}^{1/k}(k)$ be the asymptotic growth rate for the Fibonacci sequence of order $r - 1$. Let G be a hypergraph over n nodes with cn edges generated according to the following random process. The vertices of G are partitioned into r subsets of equal size, and the edges are generated at random subject to the constraint that each edge contains exactly one vertex from each set.*

With probability $1 - o(1)$, the peeling process for the k -core in G that uses r subrounds in each round terminates after $\frac{1}{r \log \phi_{r-1} + \log(k-1)} \log \log n + O(1)$ rounds when $c < c_{k,r}^$.*

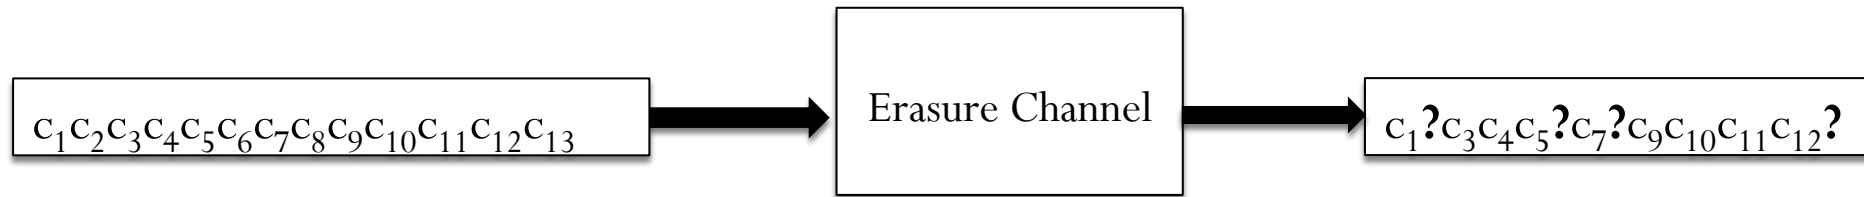
Summary: use of r subtables increase constant factor in front of the $\log \log n$, but by much less than a factor of r .

Conclusion

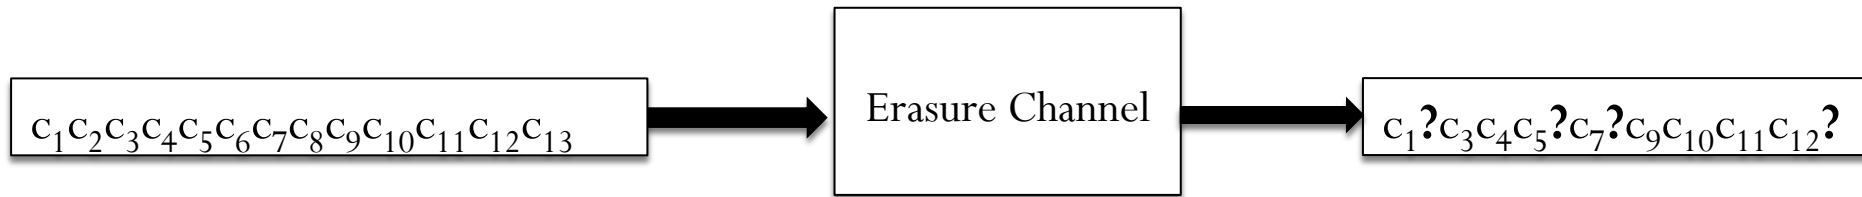
- Peeling gives simple, fast greedy algorithms.
 - Usually linear or quasi-linear total work.
- Particularly well suited for parallelization.
 - Especially when aiming for an empty k -core.
- Implementation leads to interesting variation in the analysis.
 - Subrounds.
- Can analyze dependence on “gap” to the threshold.

Thank you!

Example 1: LDPC Codes for Erasure Channels

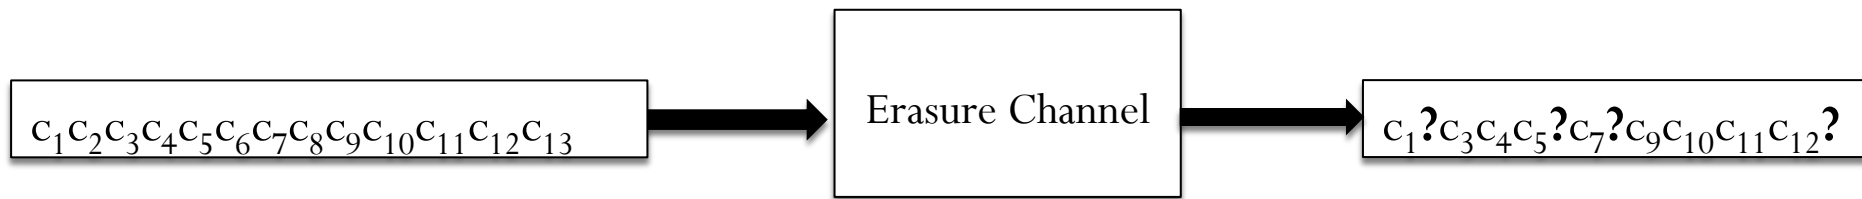


Example 1: LDPC Codes for Erasure Channels

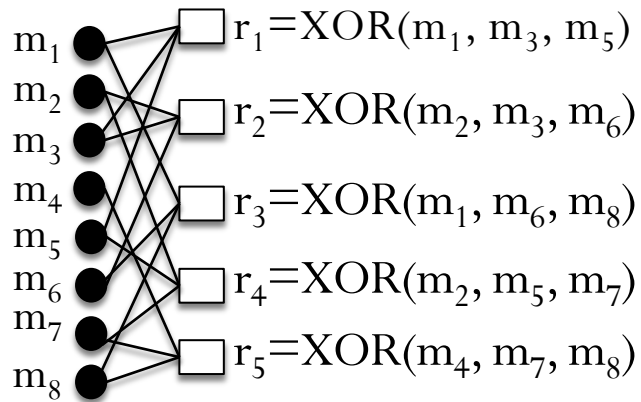


How does an LDPC code encode an 8-bit message $m_1 m_2 m_3 m_4 m_5 m_6 m_7 m_8$?

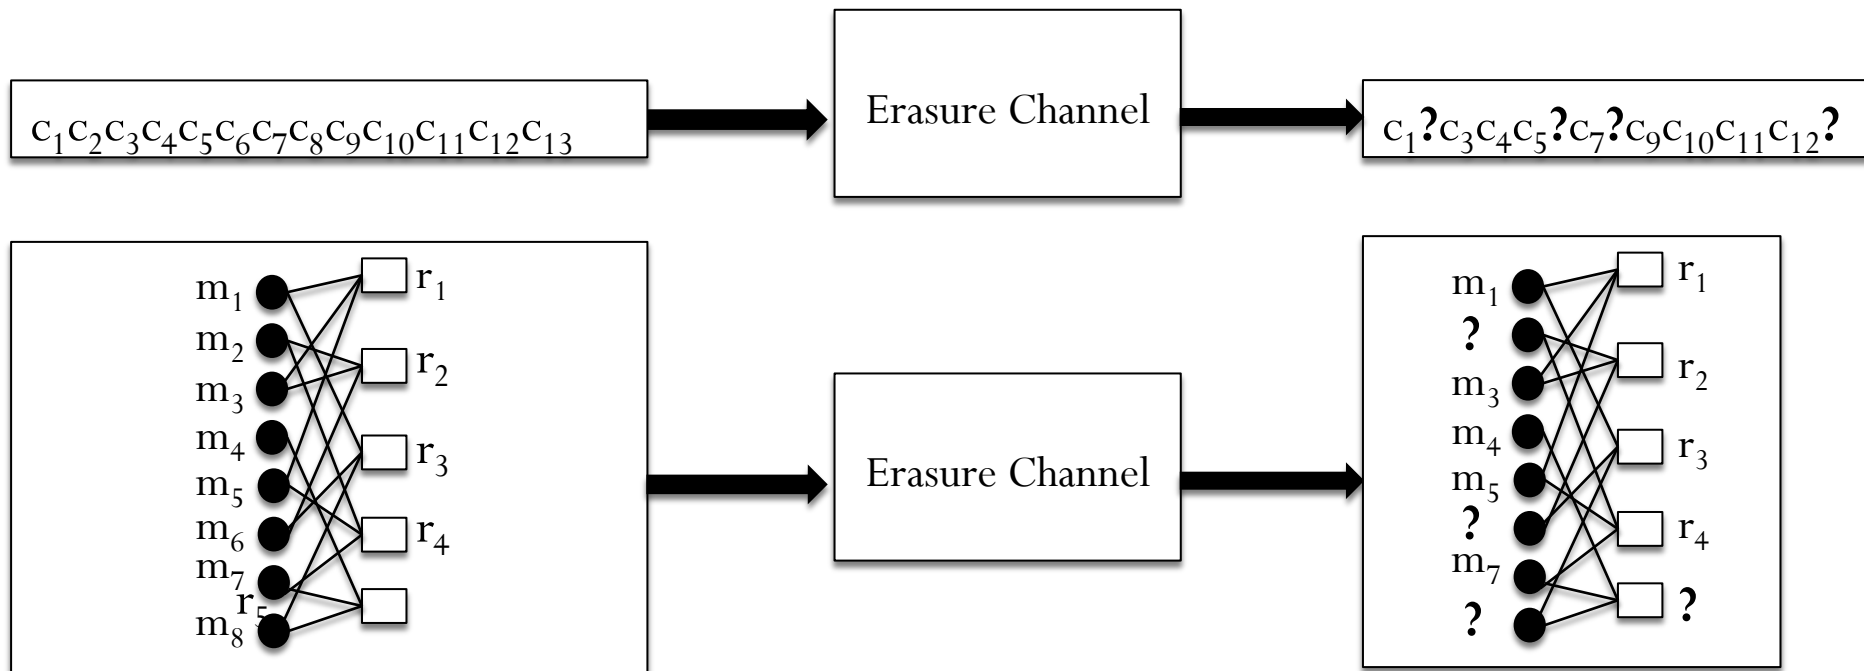
Example 1: LDPC Codes for Erasure Channels



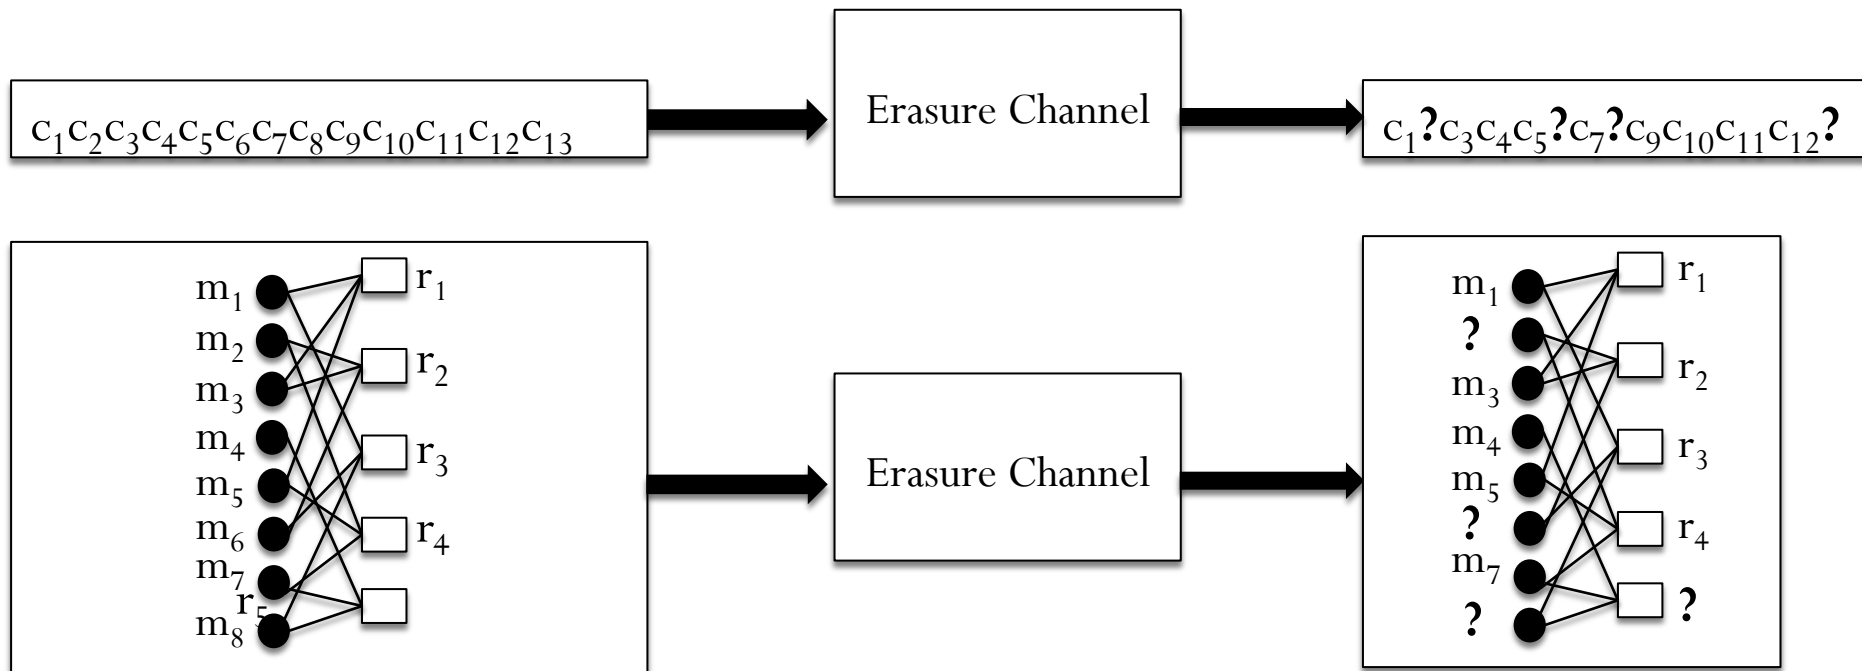
How does an LDPC code encode an 8-bit message $m_1 m_2 m_3 m_4 m_5 m_6 m_7 m_8$?



Example 1: LDPC Codes for Erasure Channels



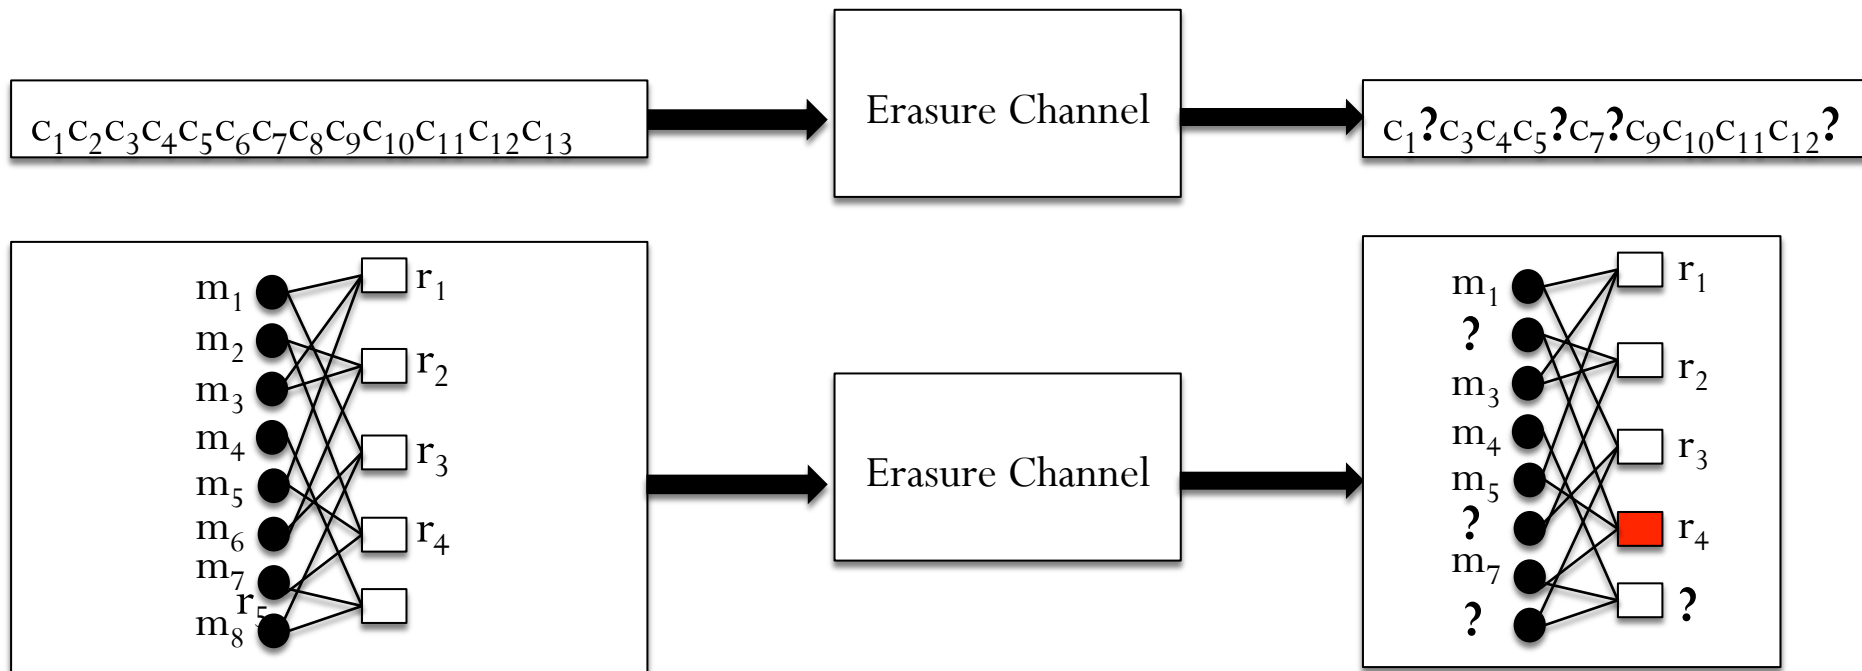
Example 1: LDPC Codes for Erasure Channels



Decoding Algorithm:

While there exists an un-erased a parity-check bit with exactly one un-erased neighbor:
Recover the neighbor

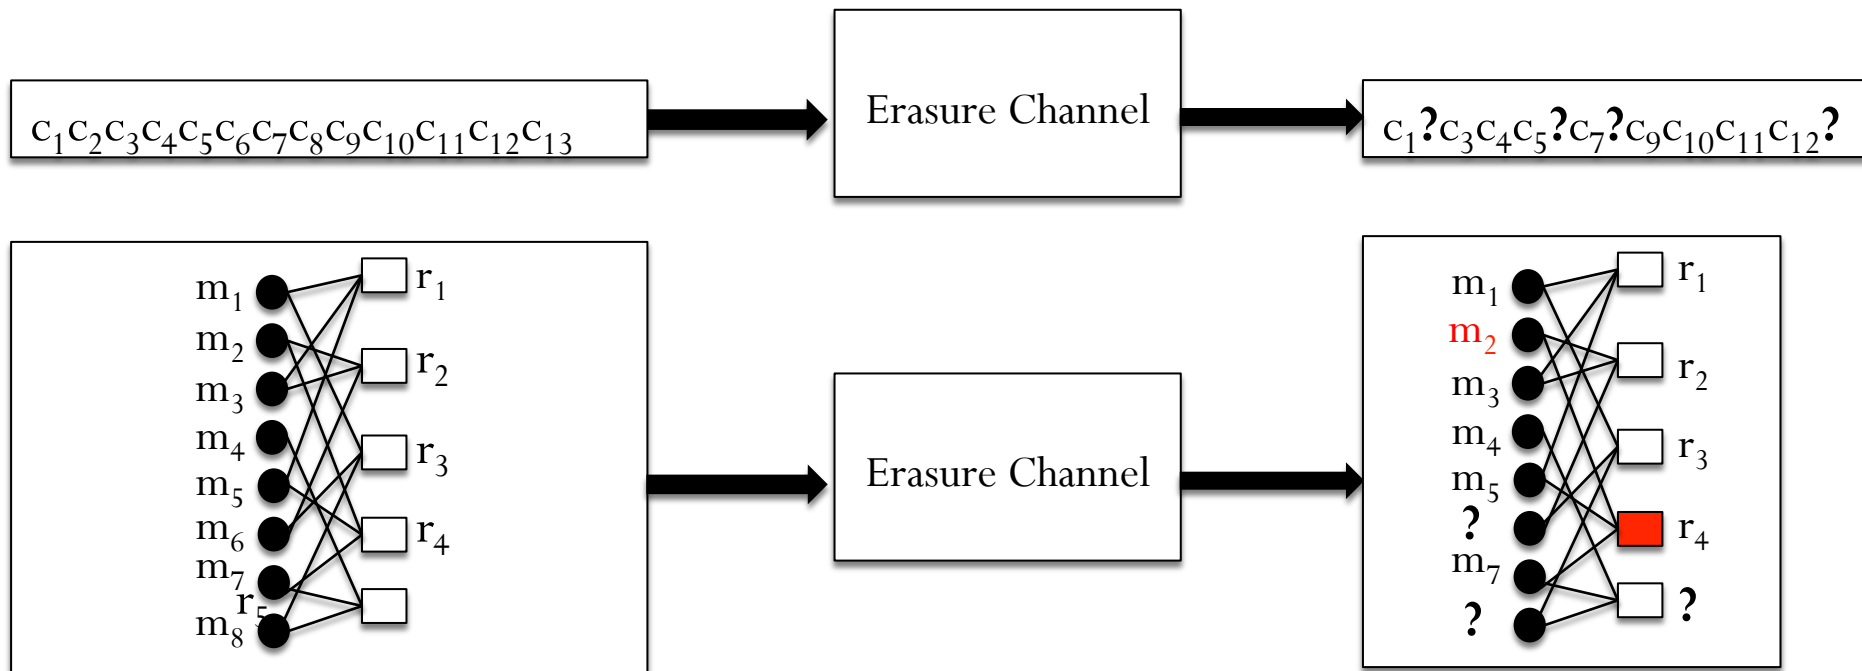
Example 1: LDPC Codes for Erasure Channels



Decoding Algorithm:

While there exists an un-erased a parity-check bit with exactly one un-erased neighbor:
Recover the neighbor

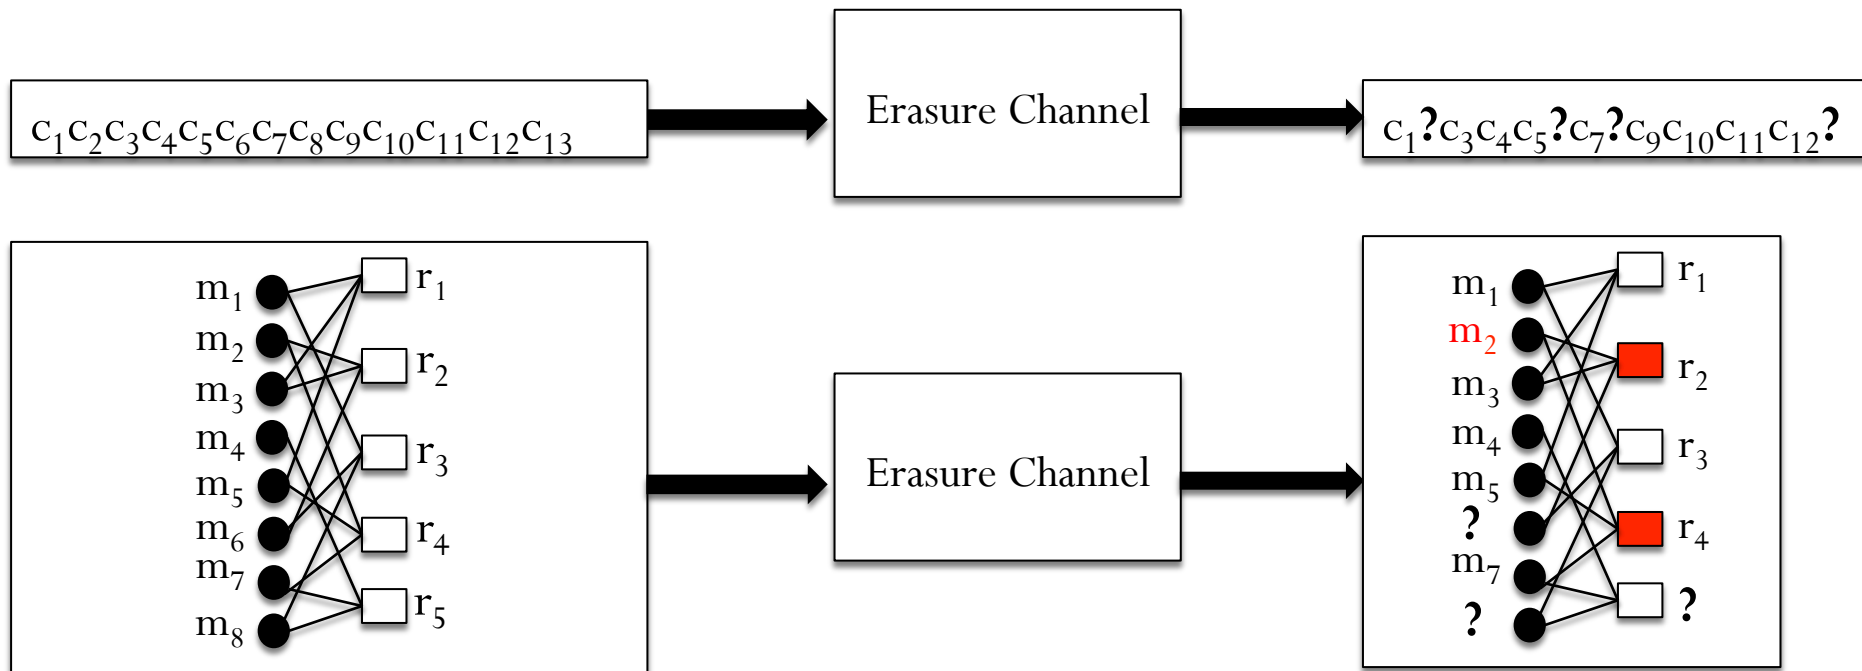
Example 1: LDPC Codes for Erasure Channels



Decoding Algorithm:

While there exists an un-erased a parity-check bit with exactly one un-erased neighbor:
Recover the neighbor

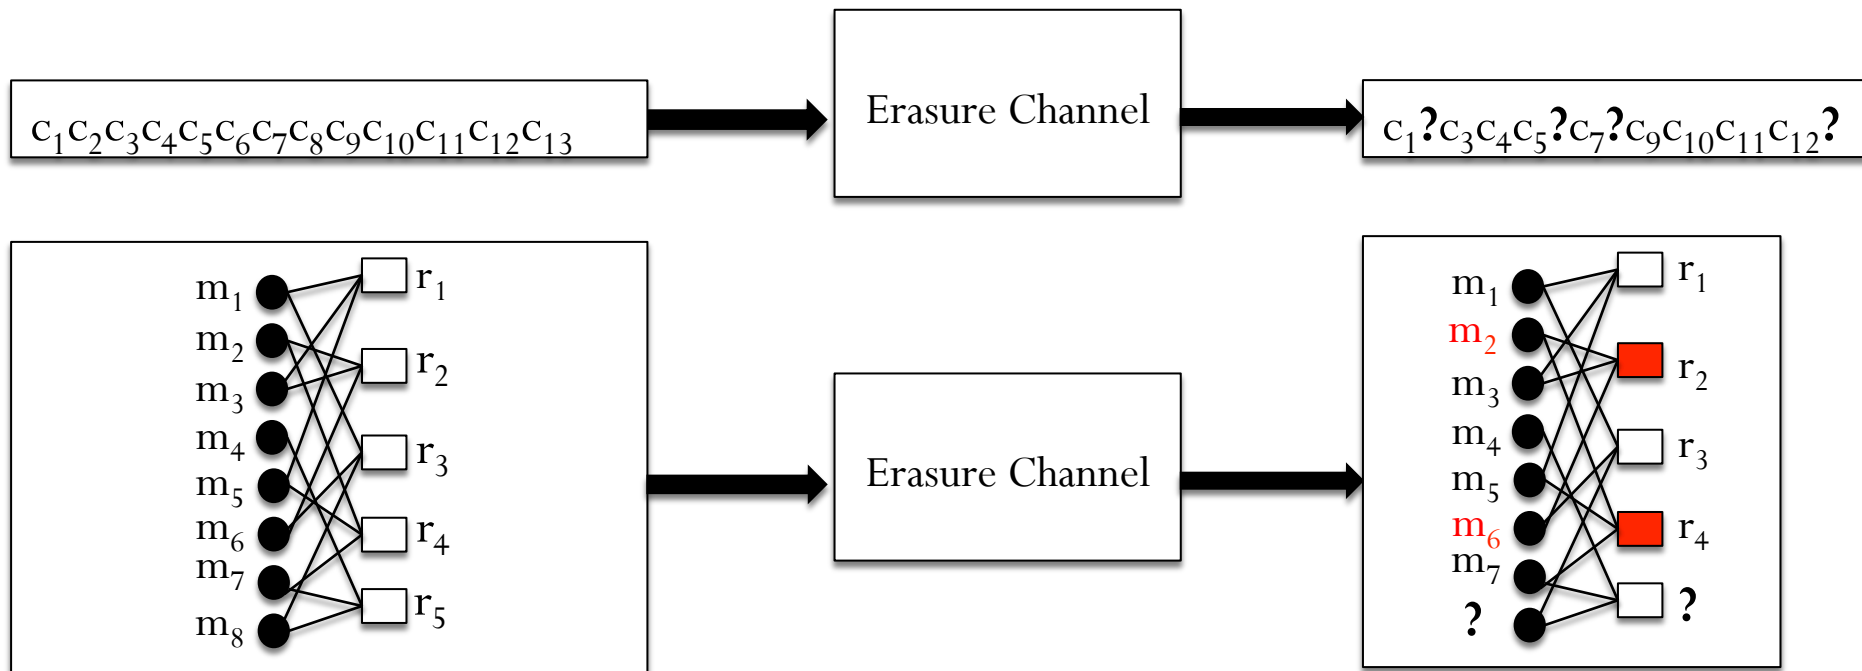
Example 1: LDPC Codes for Erasure Channels



Decoding Algorithm:

While there exists an un-erased a parity-check bit with exactly one un-erased neighbor:
Recover the neighbor

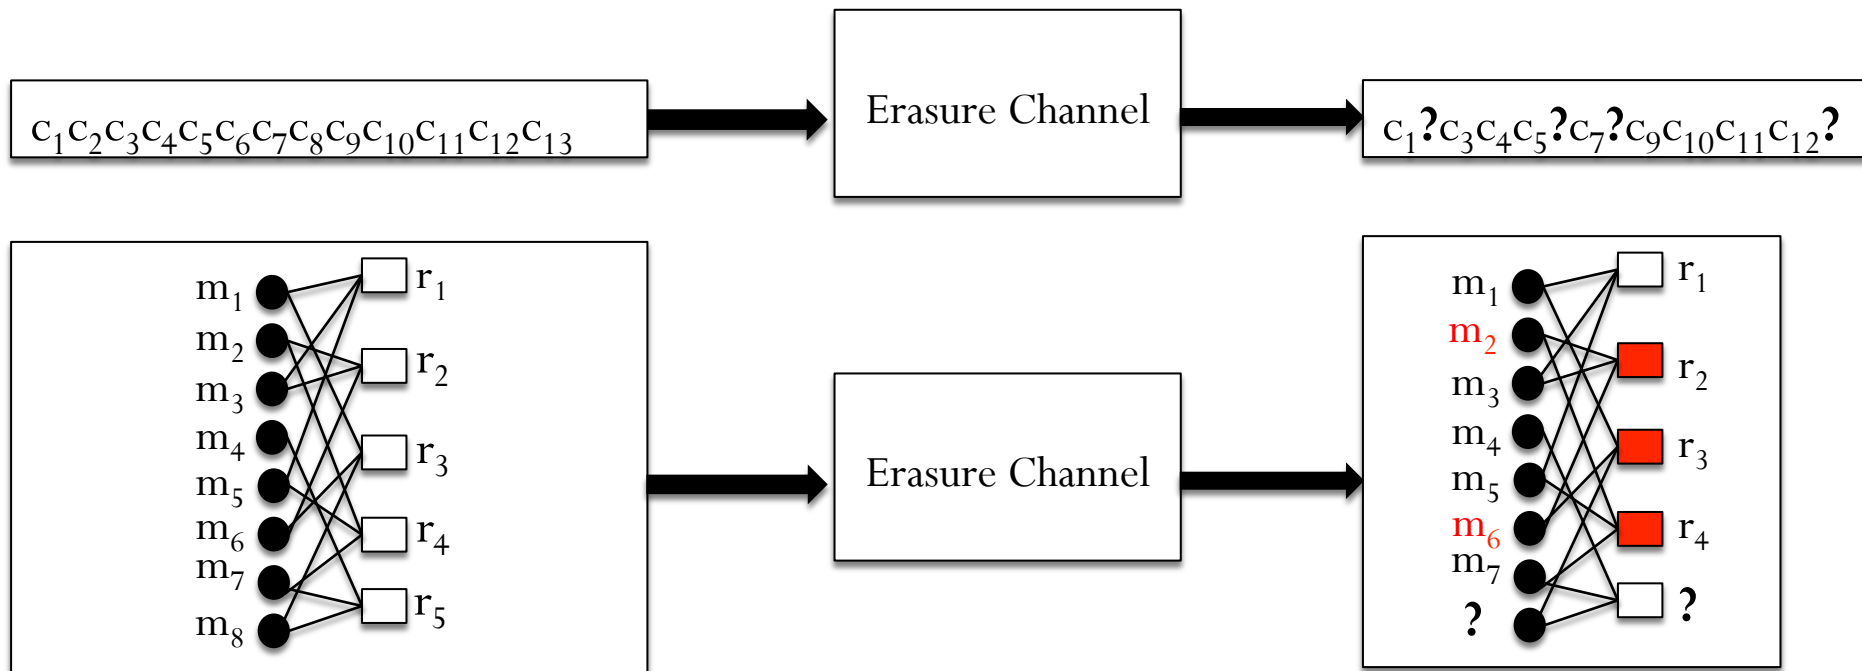
Example 1: LDPC Codes for Erasure Channels



Decoding Algorithm:

While there exists an un-erased a parity-check bit with exactly one un-erased neighbor:
Recover the neighbor

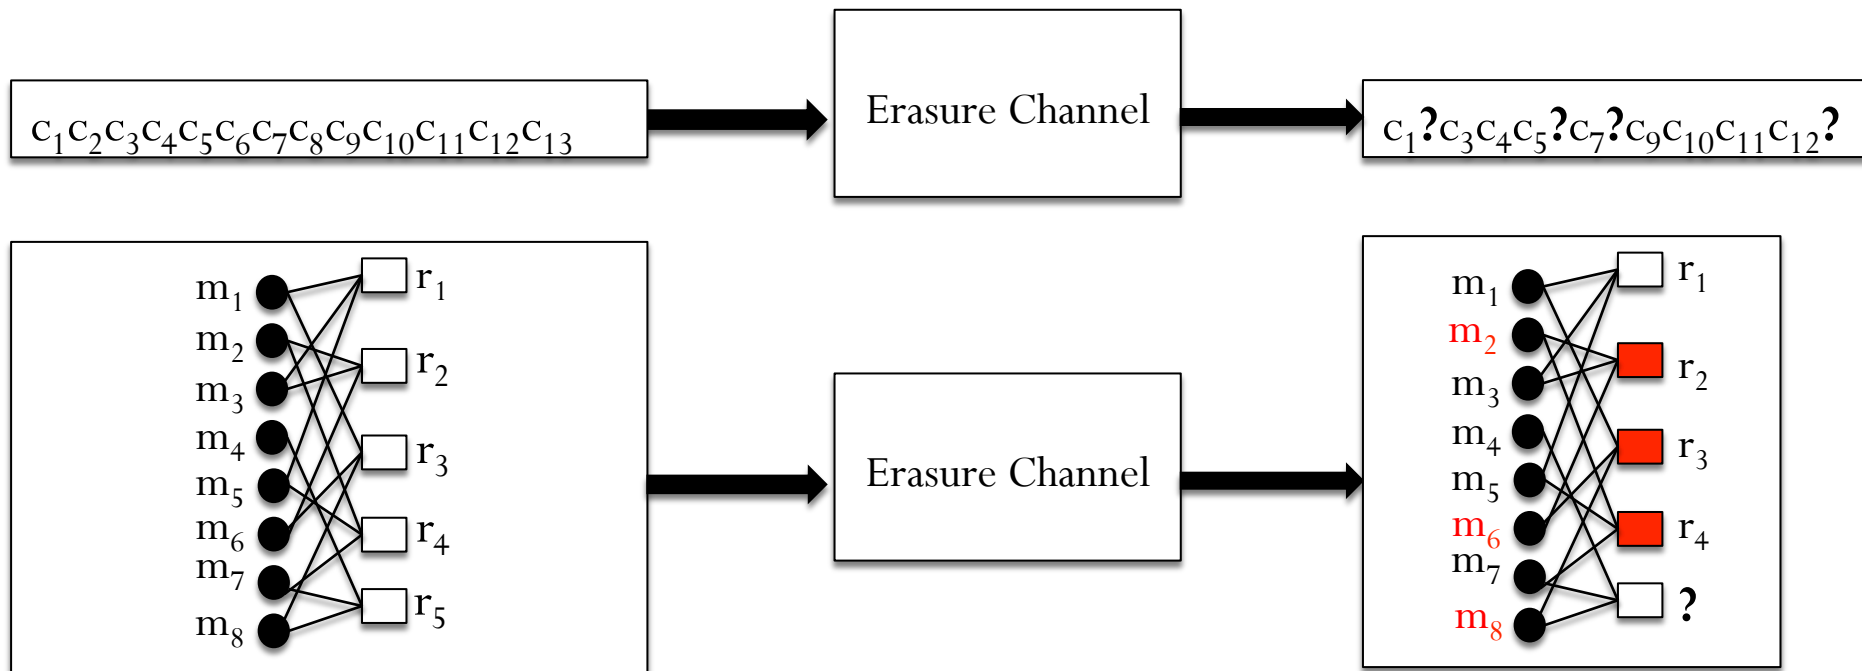
Example 1: LDPC Codes for Erasure Channels



Decoding Algorithm:

While there exists an un-erased a parity-check bit with exactly one un-erased neighbor:
Recover the neighbor

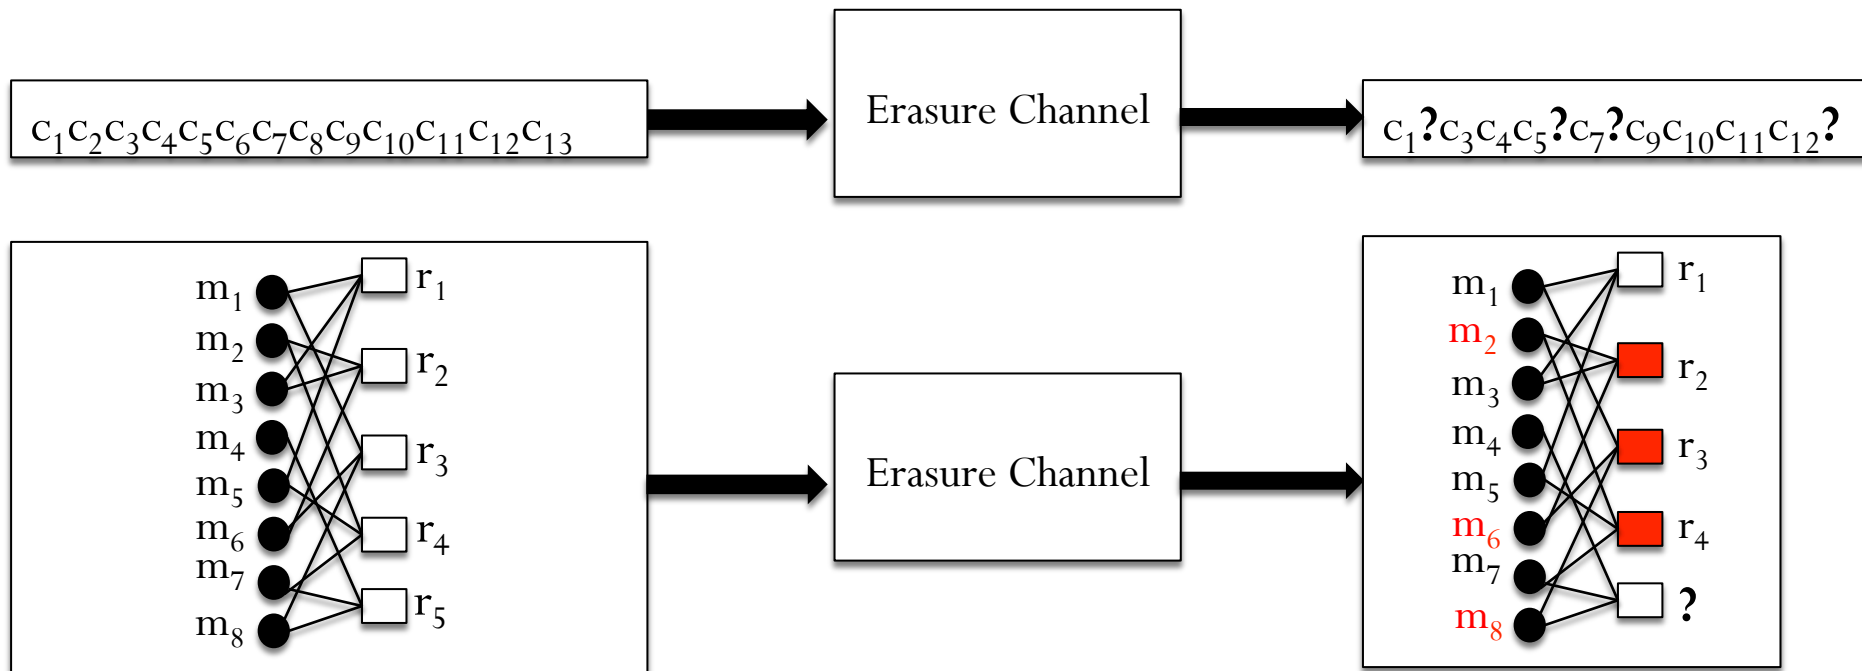
Example 1: LDPC Codes for Erasure Channels



Decoding Algorithm:

While there exists an un-erased a parity-check bit with exactly one un-erased neighbor:
Recover the neighbor

Example 1: LDPC Codes for Erasure Channels



- Decoding \longleftrightarrow peeling to 2-core on the hypergraph G where:
 - Parity-check bits \longleftrightarrow vertices of G ,
 - Erased message bits \longleftrightarrow hyperedges of G .
- Yields capacity-achieving codes with linear encoding and decoding time [Luby, Mitzenmacher, Shokrollahi, Spielman]